

Analysis of Algorithms, I

CSOR W4231

Eleni Drinea
Computer Science Department

Columbia University

Depth-first search, topological sorting

Outline

1 Recap

2 Applications of BFS

- Testing bipartiteness

3 Depth-first search (DFS)

4 Applications of DFS

- Cycle detection
- Topological sorting

Today

1 Recap

2 Applications of BFS

- Testing bipartiteness

3 Depth-first search (DFS)

4 Applications of DFS

- Cycle detection
- Topological sorting

Review of the last lecture

- ▶ Graphs (directed, undirected, weighted, unweighted)
 - ▶ Notation: $G = (V, E)$, $|V| = n$, $|E| = m$
- ▶ Representing graphs
 1. Adjacency matrix
 2. Adjacency list
- ▶ Trees, bipartite graphs, the degree theorem
- ▶ Linear graph algorithms
- ▶ Breadth-first search (BFS)

Claim 1.

Let T be a BFS tree, let x and y be nodes in T belonging to layers L_i and L_j respectively, and let (x, y) be an edge in G . Then i and j differ by at most 1.

An algorithm for s - t connectivity: breadth-first search

Breadth-first search ($\text{BFS}(G, s)$): explore G starting from s **outward in all possible directions**, adding reachable nodes one **layer** at a time.

- ▶ First add all nodes that are joined by an edge to s : these nodes form the first layer.
If G is unweighted, these are the nodes at distance 1 from s .
- ▶ Then add all nodes that are joined by an edge to a node in the first layer: these nodes form the second layer.
If G is unweighted, these are the nodes at distance 2 from s .
- ▶ And so on and so forth.

Today

1 Recap

2 Applications of BFS

- Testing bipartiteness

3 Depth-first search (DFS)

4 Applications of DFS

- Cycle detection
- Topological sorting

Testing bipartiteness & graph 2-colorability

Testing bipartiteness

- ▶ **Input:** a graph $G = (V, E)$
- ▶ **Output:** **yes** if G is bipartite, **no** otherwise

Equivalent problem (*why?*)

- ▶ **Input:** a graph $G = (V, E)$
- ▶ **Output:** **yes** if and only if we can color all the vertices in G using at most 2 colors –say red and white– so that no edge has two endpoints with the same color.

Why wouldn't we be able to 2-color a graph?

Fact: If a graph contains an odd-length cycle, then it is not 2-colorable.

So a **necessary** condition for a graph to be 2-colorable is that it does not contain odd-length cycles.

*Is this condition also **sufficient**, that is, if a graph does not contain odd-length cycles, then is it 2-colorable?*

In other words, are odd cycles the only obstacle to bipartiteness?

Algorithm for 2-colorability

BFS provides a natural way to 2-color a graph $G = (V, E)$:

- ▶ Start BFS from any vertex; color it red.
- ▶ Color white all nodes in the first layer L_1 of the BFS tree. If there is an edge between two nodes in L_1 , output **no** and stop.
- ▶ Otherwise, continue from layer L_1 , coloring red the vertices in even layers and white in odd layers.
- ▶ If BFS terminates and all nodes in V have been explored (hence 2-colored), output **yes**.

Analyzing the algorithm

Upon termination of the algorithm

- ▶ either we successfully 2-colored all vertices and output **yes**, that is, declared the graph bipartite;
- ▶ or we stopped at some level because there was an edge between two vertices of that level and output **no**; in this case, we declared the graph non-bipartite.

This algorithm is **efficient**. *Is it a **correct** algorithm for 2-colorability?*

Showing correctness

To prove correctness, we must show the following statement.

If our algorithm outputs

1. **yes**, then the 2-coloring it returns is a valid 2-coloring of G ;
2. **no**, then indeed G cannot be 2-colored by **any** algorithm (e.g., because it contains an odd-length cycle).

The next claim proves that this is indeed the case by examining the possible outputs of our algorithm. Note that the output depends solely on whether *there is an edge in G between two nodes in the same BFS layer*.

Correctness of algorithm for 2-colorability

Claim 2.

Let G be a connected graph, and let L_1, L_2, \dots be the layers produced by BFS starting at node s . Then exactly one of the following is true.

1. *There is no edge in G joining two nodes in the same BFS layer.*
Then G is bipartite and has no odd length cycles.
2. *There is an edge in G joining two nodes in the same BFS layer.*
Then G contains an odd length cycle, hence is not bipartite.

Corollary 1.

A graph is bipartite if and only if it contains no odd length cycle.

Proof of Claim 2, part 1

1. **Assume** that no edge in G joins two nodes of the same layer of the BFS tree.

By Claim 1, all edges in G not belonging to the BFS tree are

- ▶ either edges between nodes in the same layer;
- ▶ or edges between nodes in adjacent layers.

Our assumption implies that all edges of G not appearing in the BFS tree are between nodes in adjacent layers.

Since our coloring procedure gives such nodes different colors, the whole graph can be 2-colored, hence it is bipartite.

Proof of Claim 2, part 2

2. **Assume** that there is an edge $(u, v) \in E$ between two nodes u and v on the same layer.

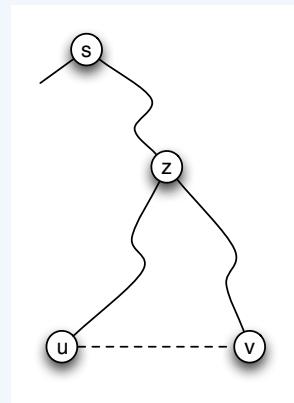
Obviously G is not 2-colorable by our algorithm: both endpoints of edge (u, v) are assigned the same color.

Our algorithm returns **no**, hence declares G non-bipartite.

*Can we show existence of an odd-length cycle and prove that G indeed is not 2-colorable by **any** algorithm?*

Proof of correctness, part 2

- ▶ Let u, v appear at layer L_j and edge $(u, v) \in E$.
- ▶ Let z be the common ancestor at max depth of u and v in the BFS tree (z might be s). Suppose z appears at layer L_i with $i < j$.
- ▶ Consider the following path in G : from z to u follow edges of the BFS tree, then to v via edge (u, v) and back to z following edges of the BFS tree. This is a cycle starting and ending at z , consisting of $(j - i) + 1 + (j - i) = 2(j - i) + 1$ edges, hence of odd length.



Today

1 Recap

2 Applications of BFS

- Testing bipartiteness

3 Depth-first search (DFS)

4 Applications of DFS

- Cycle detection
- Topological sorting

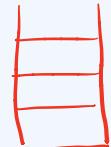
Finding your way in a maze

Application for DFS?

Use it to find your way in a maze
of interconnected rooms.

- chalk \rightarrow explored []

- string \rightarrow stack



Depth-first search (DFS): starting from a vertex s , explore the graph as deeply as possible, then **backtrack**

1. Try the first edge out of s , towards some node v .
2. Continue from v until you reach a **dead end**, that is a node whose neighbors have all been explored.
3. **Backtrack** to the first node with an unexplored neighbor and repeat 2.

Remark: DFS answers $s-t$ connectivity –

Because it explores all paths that start from s .

Use **chalk** to mark rooms that have already been explored, and use **string** to backtrack.

DFS vs BFS

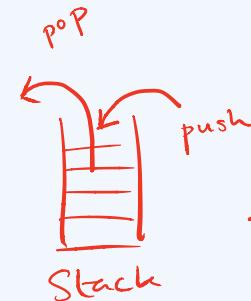
Similarities (btwn DFS and BFS)

- ▶ Linear-time algorithms that essentially can be used to perform the same tasks

Differences (btwn DFS and BFS)

- ▶ DFS is more *impulsive*: when it *discovers* an *unexplored* node, it moves on to exploring it right away; BFS defers exploring until all nodes in the layer have been discovered.
- ▶ DFS is naturally recursive and implemented using a **stack**.

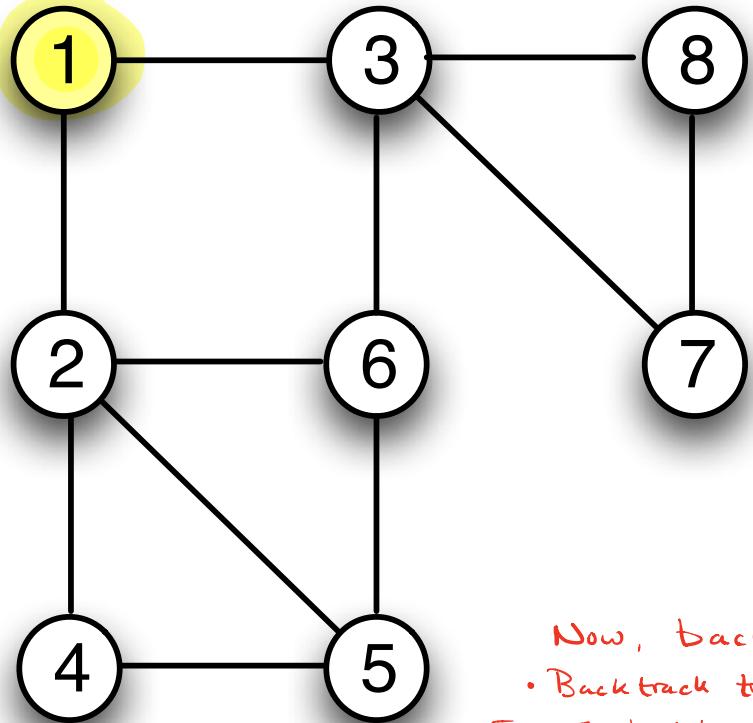
- ▶ A stack is a LIFO (Last-In First-Out) data structure implemented as a doubly linked list: **insert** (**push**)/**extract** (**pop**) the top element requires $O(1)$ time.



- Insert elements to the top of the stack (**push**)
- Extract the top element (**pop**). Always act on the top of stack.

An undirected graph G_1

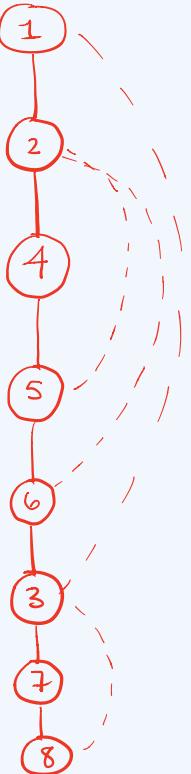
START



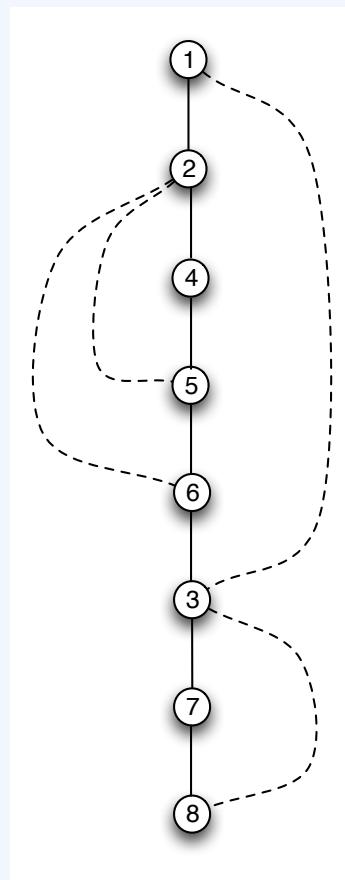
Output of
DFS is a
tree

Now, back track.

- Backtrack to 7
- From 7, backtrack to 3
- 3 has no unexplored neighbors, backtrack to 6.
- From 6, No unexplored neighbors, backtrack to node 5.
- Node 5 has no unexplored neighbors. Backtrack to 4.
- ... to 2 ... to 1
- DFS terminates.

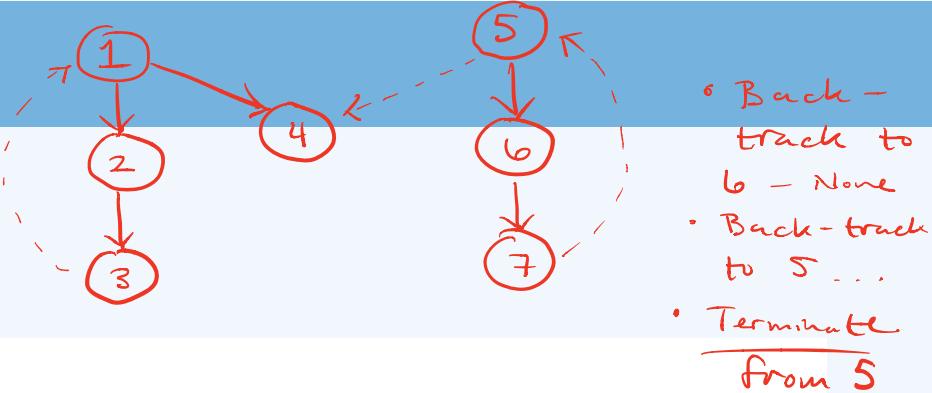
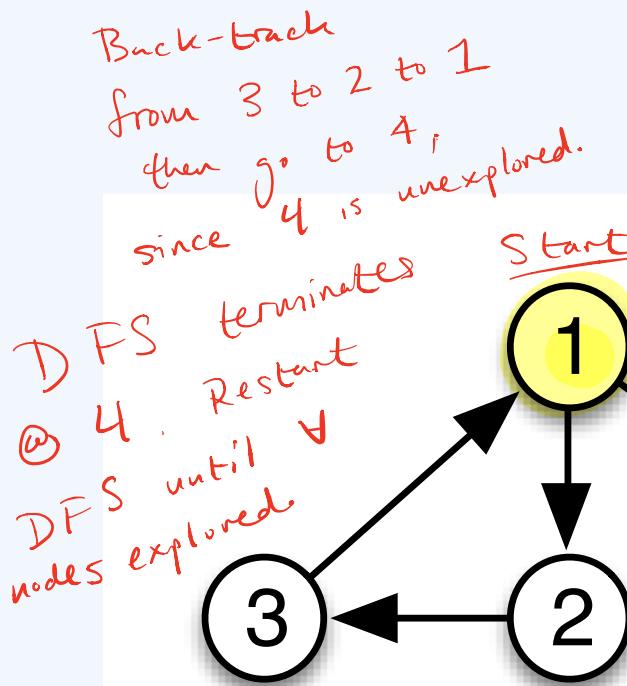


The DFS tree for the undirected graph G_1



Dashed edges belong to the graph but not to the DFS tree. Ties are broken by considering nodes by increasing index.

A directed graph G

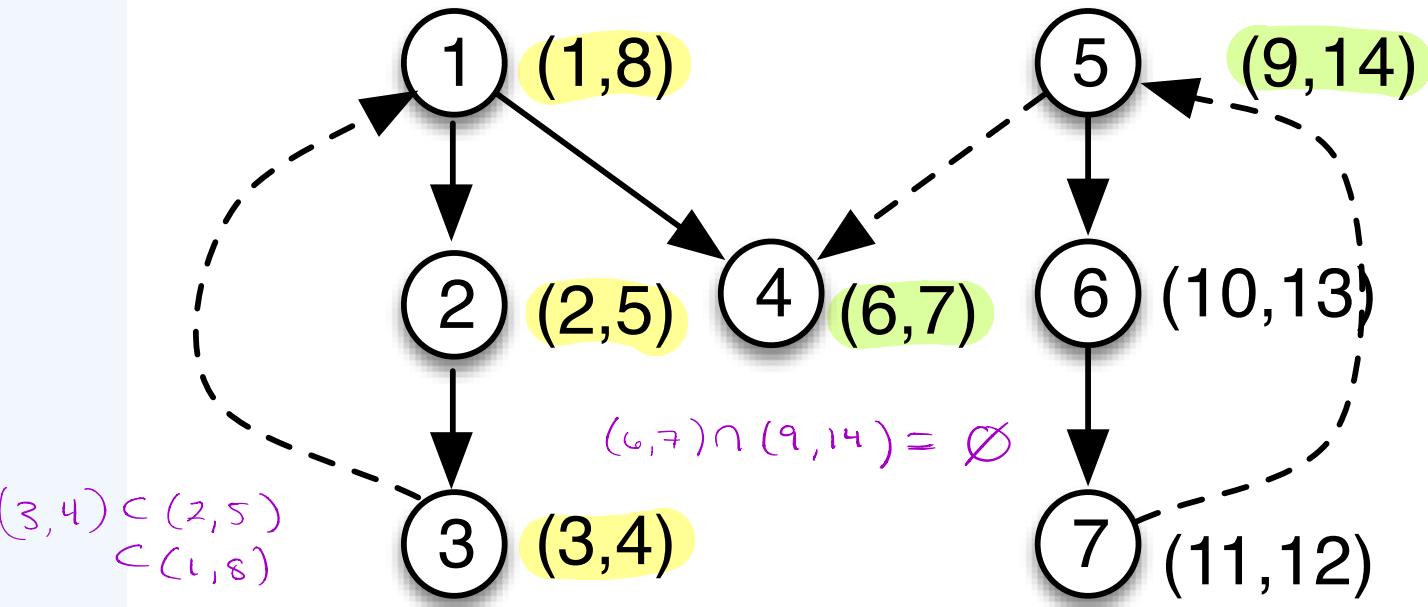


Restart DFS until all nodes in the graph have been explored.

• Output of DFS in this case is a forest of trees.

The DFS forest for the directed graph G

Notice the relationship of the time intervals



Two possible relationships between time intervals:
(i) Containment (ii) Disjoint

Dashed edges belong to G but not to the trees in the DFS forest.
($start, finish$) intervals appear to the right of every node.

Pseudocode for DFS exploration of the entire graph

```

DFS( $G = (V, E)$ )
 $O(n) \left\{ \begin{array}{l} \text{for } u \in V \text{ do} \\ \quad explored[u] = 0 \\ \text{end for} \\ \text{for } u \in V \text{ do} \\ \quad \text{if } explored[u] == 0 \text{ then Search}(u) \\ \quad \text{end if} \\ \text{end for} \end{array} \right.$ 

```

Search(u)

previsit(u) $\rightarrow O(1)$

$explored[u] = 1 \rightarrow O(1)$

for $(u, v) \in E$ do

if $explored[v] == 0$ then Search(v)

end if

end for

postvisit(u) $\rightarrow O(1)$

$$\text{Total Time} = O(n + m)$$

For fixed u ,
How many times do
you call $\text{Search}(u)$? ONCE.

Let $T(u) =$ time for $\text{Search}(u)$

EXCLUDING the time
spent on recursive calls
inside $\text{Search}(u)$.

$$* T(1) + T(2) + \dots + T(n)$$

$$= \sum_{u \in V} T(u) = c \cdot \sum_{u \in V} \deg(u)$$

$$= c \cdot 2m$$

$$= O(m)$$

For fixed u :

time for comparison
($explored[v] == 0$)

$$T(u) = c' + O(1) \cdot \deg(u) = c \cdot \deg(u)$$

Running time for DFS if previsit, postvisit take $O(1)$ time?

Directed graphs: classification of edges

Remember these classifications.

Important for structures of the DFS graph in our algorithms.

Graph edges that do not belong to the DFS tree(s) may be

1. **forward**: from a vertex to a *descendant* (other than a *child*)

2. **back**: from a vertex to an *ancestor*

Examples: edges $(3, 1), (7, 5)$ in G

must be in the
same DFS tree.

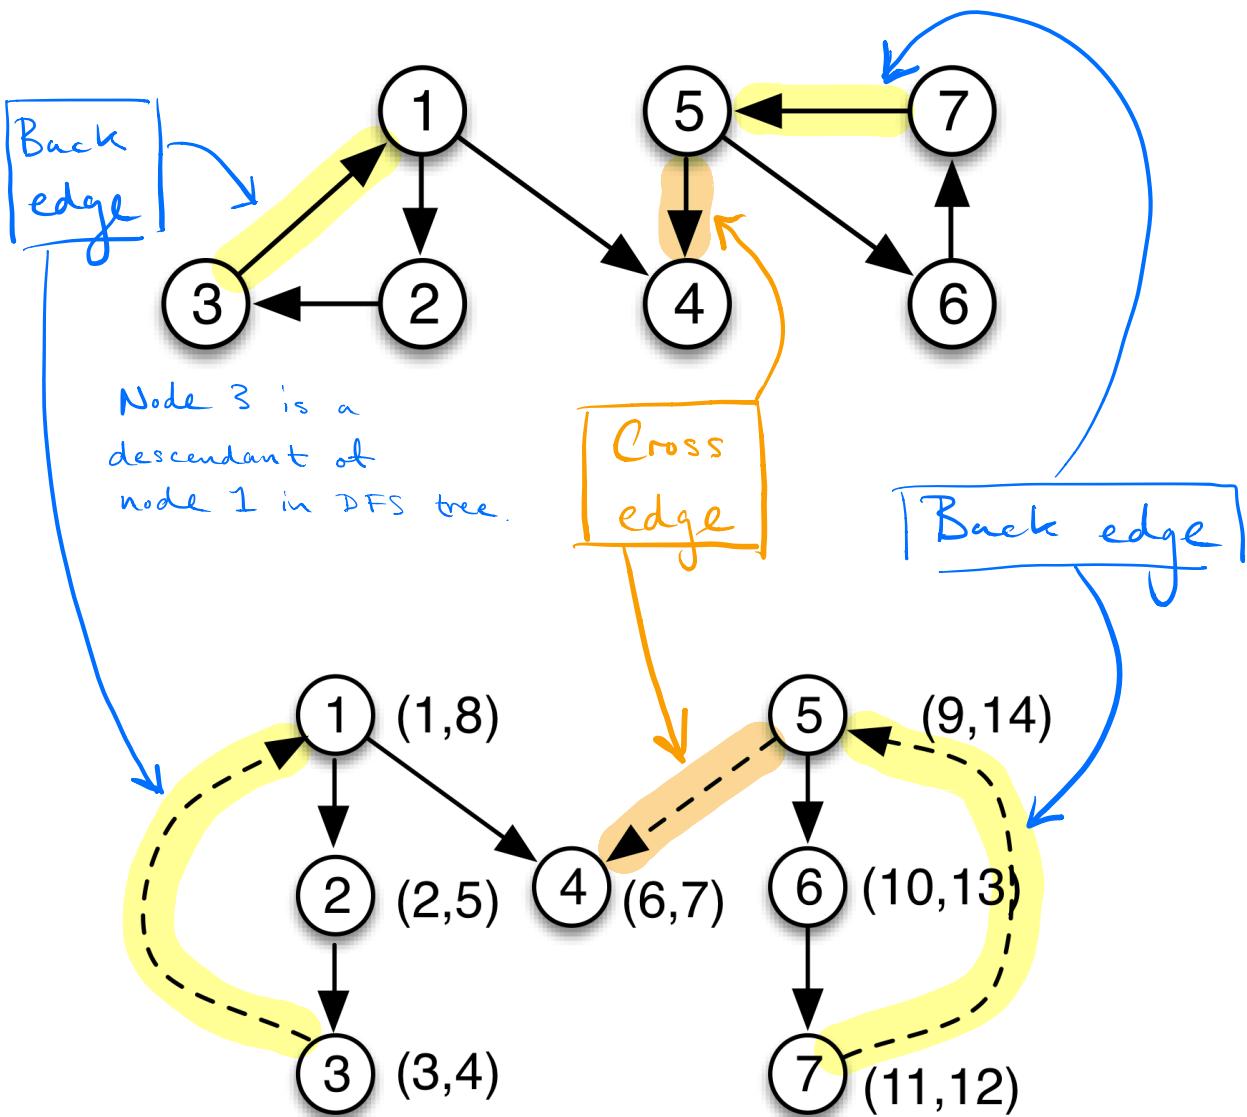
3. **cross**: from right to left (no ancestral relation), that is

- ▶ from tree to tree (example: edge $(5, 4)$ in G)

- ▶ between nodes in the same tree but on different branches

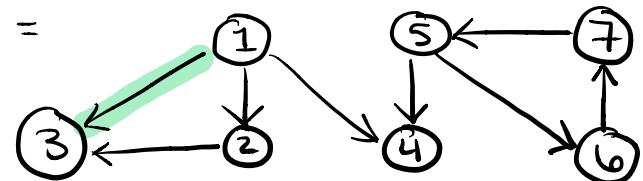
Ancestral
relationships
w.r.t. the
DFS tree
output.

④ and ⑤ belong to different
trees in the DFS output,
but don't have to, to be a
cross edge.



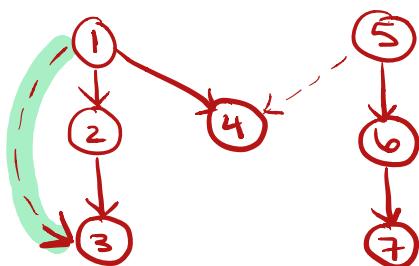
Example of forward edge:

$G =$



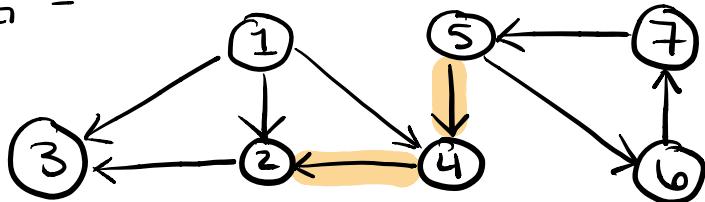
$(1, 3)$ is a forward edge here.

DFS output:

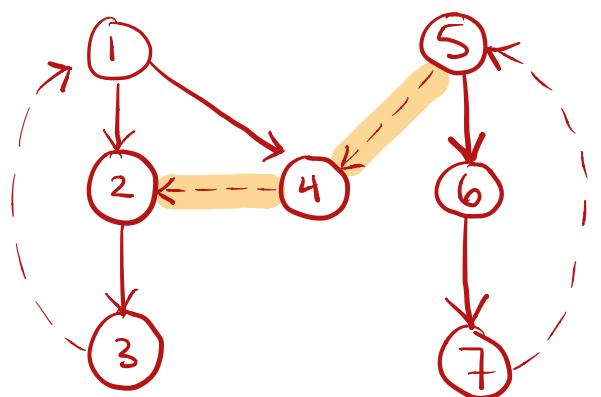


Cross edge example

$G =$



DFS output :



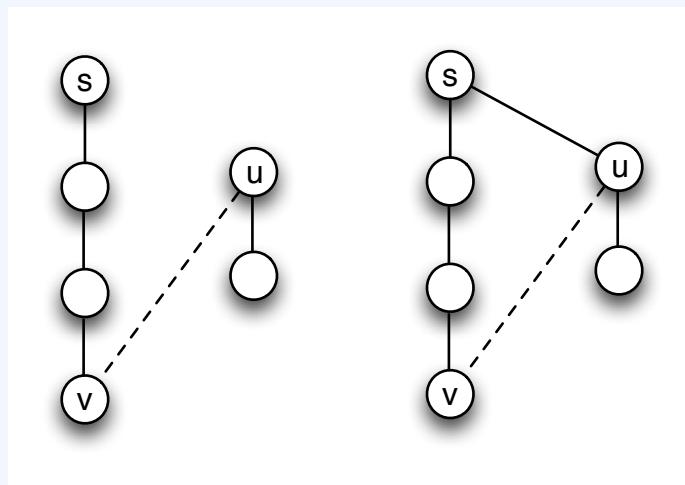
Cross edges :

(4, 2) and (5, 4)

There is no ancestral relationship between nodes 4 and 2. That is, they are not on the same path in the output of DFS. Likewise for nodes 5 and 4.

Undirected graphs: classification of edges

Cross and *forward* edges do not exist in undirected graphs.

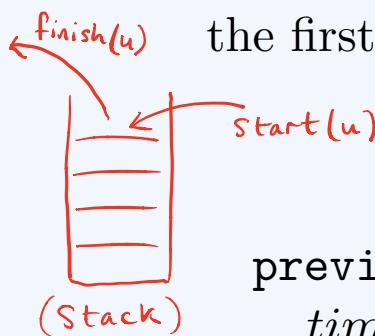


In undirected graphs, DFS only yields *back* and *tree* edges.

Time intervals for vertices

Subroutines $\text{previsit}(u)$, $\text{postvisit}(u)$ may be used to maintain a notion of **time**:

- ▶ In $\text{DFS}(G)$, initialize a counter *time* to 0.
- ▶ Increment the counter by 1 every time $\text{previsit}(u)$, $\text{postvisit}(u)$ are accessed.
- ▶ Store the times $\text{start}(u)$ and $\text{finish}(u)$ corresponding to the first and last time u was visited during $\text{DFS}(G)$.



$\text{previsit}(u)$

$$\text{time} = \text{time} + 1$$

$$\text{start}(u) = \text{time}$$

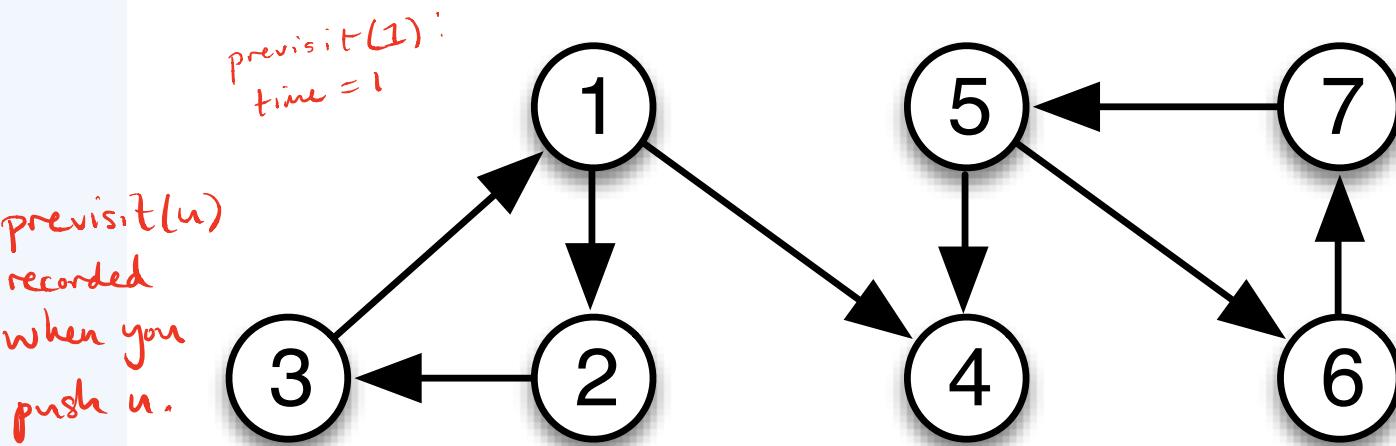
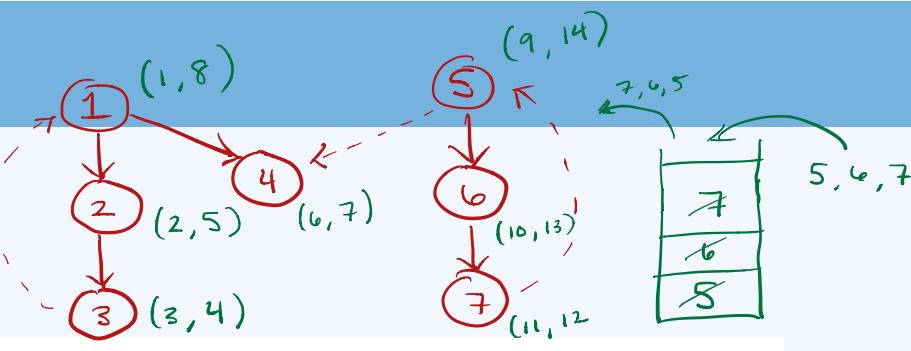
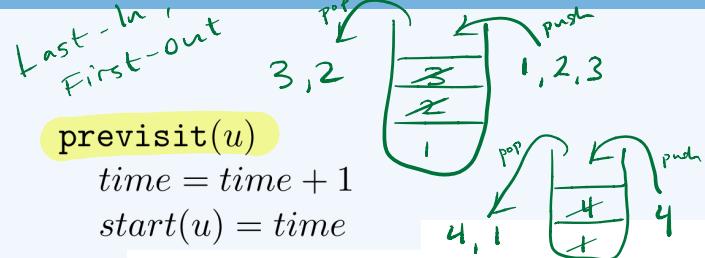
$\text{postvisit}(u)$

$$\text{time} = \text{time} + 1$$

$$\text{finish}(u) = \text{time}$$

The interval $(\text{start}(u), \text{finish}(u))$ tells you the interval of time during which vertex u was active in the algorithm. ↴ (see below for example)

A directed graph G

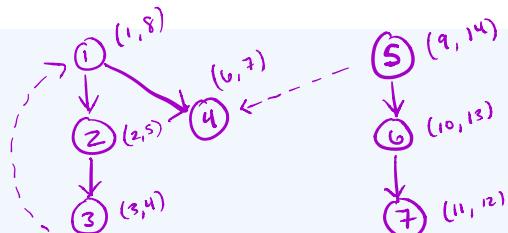


postvisit(u)

$$\text{time} = \text{time} + 1$$

$$\text{finish}(u) = \text{time}$$

postvisit(u) recorded when you pop u .



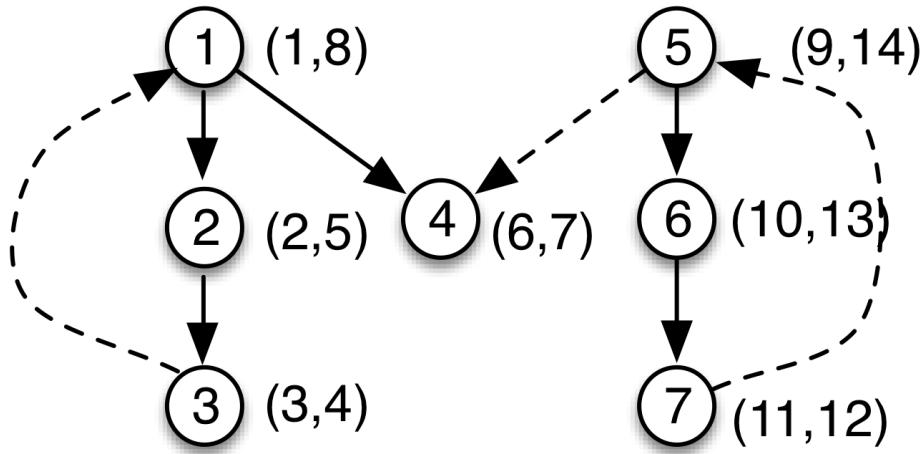
$(5, 4)$ is a cross edge
 $\Rightarrow \text{start}(v) < \text{start}(u)$
 since v already explicit.
 and $\text{finish}(v) < \text{finish}(u)$

On the time intervals of vertices u, v

If we use an explicit stack, then

- ▶ $start(u)$ is the time when u is pushed in the stack
- ▶ $finish(u)$ is the time when u is popped from the stack
(that is, all of its neighbors have been explored).

1. How do intervals $[start(u), finish(u)]$, $[start(v), finish(v)]$ relate?
Two possible relationships between the intervals:
(i) Containment (ii) Disjoint
2. What do the contents of the stack correspond to in the DFS tree (and the graph), if s was the first vertex pushed in the stack and v the last?
 - Corresponds to a path
- Interval containment corresponds to ancestral relationships
 - The interval of the descendant is fully contained in the interval of the ancestor.



- The time interval of descendants is fully contained within the interval of the parent node.
- If there is no ancestral relation, the time intervals are disjoint

Identifying back edges using time

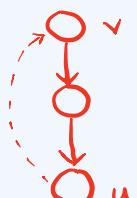
1. Intervals $[start(u), finish(u)]$ and $[start(v), finish(v)]$
 - ▶ either contain each other (u is an ancestor of v or vice versa)
 - ▶ or they are disjoint.
2. If s was the first vertex pushed in the stack and v is the last, the vertices currently in the stack form an $s-v$ path.

Claim 3 (Back edges).

Let $(u, v) \in E$. Edge (u, v) is a back edge in a DFS tree if and only if

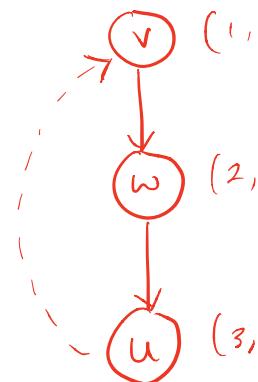
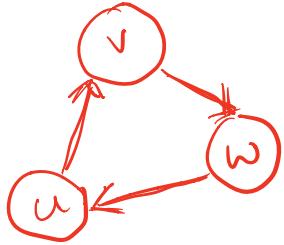
$$start(v) < start(u) < finish(u) < finish(v).$$

(e.g. interval of u is fully contained within the interval of v)



Use this fact to identify back edges in the DFS output.

$\underline{\underline{Pf}}:$ (u, v) back edge in DFS tree
 $\Rightarrow \text{start}(v) < \text{start}(u) < \text{finish}(u) < \text{finish}(v)$



If (u, v) a back-edge

then v must be an ancestor
of $u \Rightarrow v$ enters the stack
before u

$\Rightarrow \text{start}(v) < \text{start}(u)$

and clearly $\text{start}(u) < \text{finish}(u)$.

Since v is first, then by
construction of DFS alg.

v leaves last,

so $\text{finish}(u) < \text{finish}(v)$

$\Rightarrow \text{start}(v) < \text{start}(u) < \text{finish}(u) < \text{finish}(v)$

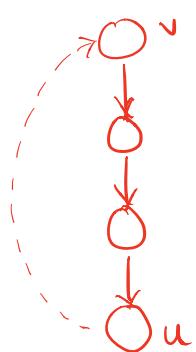
\Leftarrow Given the inequality

$$\text{start}(v) < \text{start}(u) < \text{finish}(u) < \text{finish}(v)$$

\Rightarrow Since $\text{start}(v) < \text{start}(u)$

then v must be an ancestor
of u .

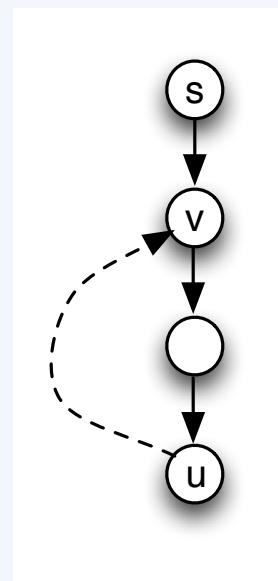
This implies there is a path from v to u



if there is an edge (u,v) in the graph then the edge (uv) in the DFS tree must be a back edge.

Proof of Claim 3 (identifying back edges)

Proof.



If (u, v) is a back edge, the claim follows.

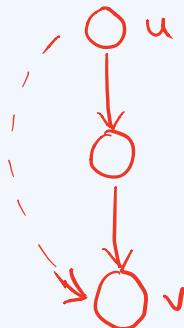
Otherwise, v was pushed in the stack before u and is still in the stack when u is pushed into it. Then there is a $v - u$ path in the DFS tree, so v is an ancestor of u and (u, v) is a back edge. \square

Identifying forward and cross edges

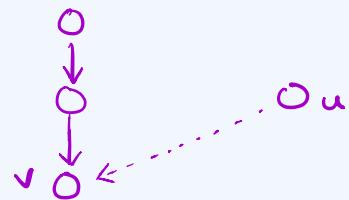
What conditions must the start and finish numbers satisfy if

1. $(u, v) \in E$ is a **forward** edge in the DFS tree? Not iff
2. $(u, v) \in E$ is a **cross** edge in the DFS tree? iff

(u, v) forward edge



(u, v) a cross edge



iff $\text{start}(v) < \text{finish}(v) < \text{start}(u) < \text{finish}(u)$

interval of v fully contained in the interval of u .

$\text{start}(u) < \text{start}(v) < \text{finish}(v) < \text{finish}(u)$

Identifying forward and cross edges

$$s(v) <$$

What conditions must the start and finish numbers satisfy if

1. $(u, v) \in E$ is a **forward** edge in the DFS tree?
2. $(u, v) \in E$ is a **cross** edge in the DFS tree?

1. If edge $(u, v) \in E$ is a forward edge, then

$$\text{start}(u) < \text{start}(v) < \text{finish}(v) < \text{finish}(u).$$

the interval of v
is fully contained within
the interval of u

2. If edge $(u, v) \in E$ is a cross edge, then

$$\text{start}(v) < \text{finish}(v) < \text{start}(u) < \text{finish}(u).$$

Not iff : Inequality implies two possible cases : (i) edge is part of the DFS tree ; (ii) edge is a forward edge.

OUTPUT OF DFS :

- For **undirected graphs**, the output of DFS is all the connected components of the input graph. If the graph is not connected, this will take the form as a forest of trees. If the graph is connected, the output will be a single, undirected tree.
- For **directed graphs**, the output of DFS is a forest of directed trees.

Today

$G = (V, E)$ has a cycle

1 Recap



DFS finds a back edge

2 Applications of BFS

- Testing bipartiteness

3 Depth-first search (DFS)

4 Applications of DFS

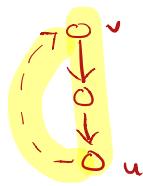
- Cycle detection
- Topological sorting

Claim 4.

$G = (V, E)$ has a cycle if and only if $\text{DFS}(G)$ yields a back edge.

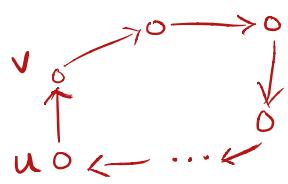
Pf :

(\Leftarrow) Trivial ;



gives
cycle
in the
graph
 G .

(\Rightarrow) Spz $G = (V, E)$ has a cycle



w.l.o.g. let v be the first node
of the cycle discovered by DFS, and
let u be the node preceding v
in the cycle.

All nodes in the cycle will appear in the
DFS sub-tree of v since there is path
from any node in the cycle to all other nodes
in the cycle (all nodes are reachable from v).

That is, since v is the first node in
the cycle, all nodes in the cycle appear
below v in the DFS output tree

Since there is a path from v to every
node in the cycle. Every node will appear, including
 u . More, since u precedes v , there is an edge
from u to v . Since v is the first node in the DFS
tree, and has already been explored, (u, v) is a
back edge. since it is an edge from a descendant
to an ancestor.

Application I: Cycle detection

Claim 4.

$G = (V, E)$ has a cycle if and only if $\text{DFS}(G)$ yields a back edge.

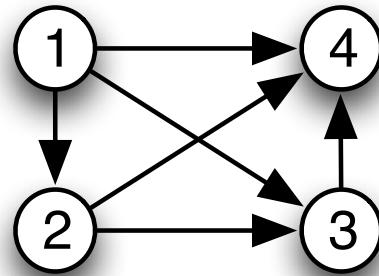
Proof.

If (u, v) is a back edge, together with the path on the DFS tree from v to u , it forms a cycle.

Conversely, suppose G has a cycle. Let v be the first vertex from the cycle discovered by $\text{DFS}(G)$. Let (u, v) be the preceding edge in the cycle. Since there is a path from v to *every* vertex in the cycle, all vertices in the cycle are now discovered **and** fully explored **before** v is popped from the stack. Hence the interval of u is contained in the interval of v . By Claim 1, (u, v) is a back edge. □

Application II: Topological sorting in DAGs

- ▶ An undirected acyclic graph has an extremely simple structure: it is a tree, hence a sparse graph ($O(n)$ edges).
- ▶ A directed acyclic graph (**DAG**) may be dense ($\Omega(n^2)$ edges): e.g., $V = \{1, \dots, n\}$, $E = \{(i, j) \text{ if } i < j\}$.



Topological sorting: motivation

Input:

- ▶ a set of tasks $\{1, 2, \dots, n\}$ that need to be performed
- ▶ a set of dependencies, each of the form (i, j) , indicating that task i must be performed before task j .

Output: a valid order in which the tasks may be performed, so that all dependencies are respected.

Example: tasks are courses and certain courses must be taken before others.

How can we model this problem using a graph? What kind of graph must arise and why?

Topological ordering: definition

Definition 2.

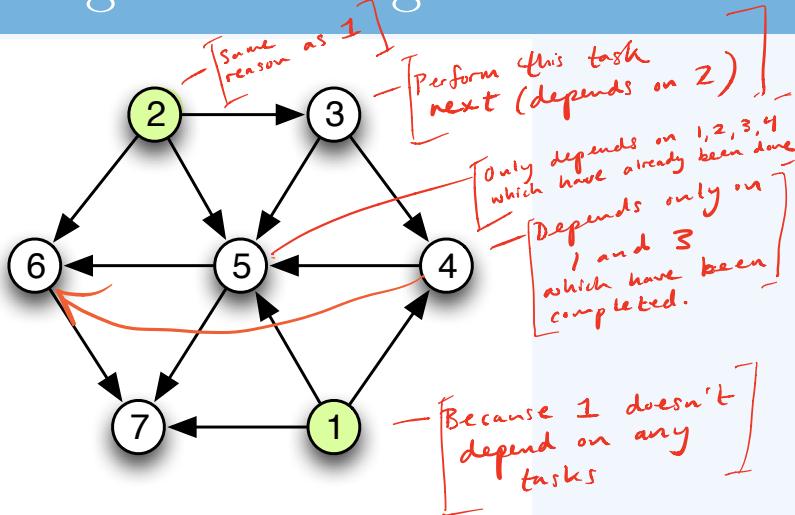
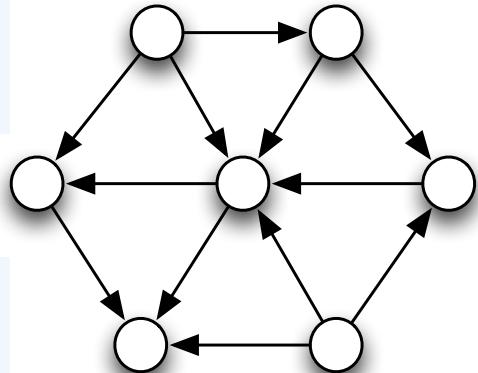
A topological ordering of G is an ordering of its nodes as $1, 2, \dots, n$ such that for every edge (i, j) , we have $i < j$.

originating end pt has an index smaller than the destination end pt.

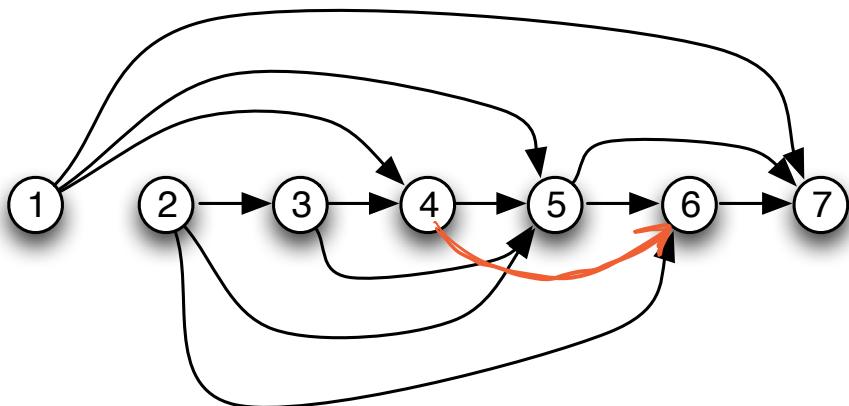
- ▶ All edges point **forward** in the topological ordering.
- ▶ It provides an order in which all tasks can be safely performed: when we try to perform task j , all tasks required to precede it have already been done.

Example of DAG and its topological sorting

DAG



If a graph has a topological ordering, then it must be a DAG.



Visual proof that this graph is a DAG:
All edges are going from left to right.

A DAG (top left), its topological sort (top right) and a drawing emphasizing the topological sort (bottom).

Topological sorting in DAGs

Claim 5.

If G has a topological ordering, then G is a DAG.

Proof: By contradiction (*exercise*).

A visual proof is provided by the linearized graph of the previous slide: vertices appear in increasing order, edges go from left to right, hence no cycles.

{ Is the converse true: does every DAG have a topological ordering? And how can we find it?

* Yes, every DAG has a topological ordering.



G has a topological ordering $\Rightarrow G$ is a DAG

Is the converse true? That is, is it true that:

G is a DAG $\Rightarrow G$ has a topological ordering?

Yes; Consider the following:



Suppose $|V| = n$,
if n had an
outgoing edge here,
the graph would
have a cycle.

- Does every node in a DAG have an outgoing edge?

| No. If every node had an outgoing edge, then the graph has a cycle which is not possible for a DAG.

This property implies :

No \Rightarrow Every DAG has ≥ 1 sink

Ex:

Every DAG has at least one sink.

- Does every node in a DAG have an incoming edge?

No; Same argument

It would otherwise have a cycle.

This property implies:

No \Rightarrow Every DAG has ≥ 1 source



Every DAG has at least one source.

To find a topological ordering for a DAG, start with a source node. We know at least one source node exists for all DAGs.

Structural properties of DAGs

In a DAG, can **every** vertex have

- ▶ *an outgoing edge?*
- ▶ *an incoming edge?*

Definition 3 (source and sink).

A **source** is a node with no incoming edges.

A **sink** is a node with no outgoing edges.

Fact 4.

Every DAG has at least one source and at least one sink.

To find a topological ordering for a DAG,
start with a source node (We know at least
one source node exists for all DAGs.). :

- (I) Label the source node first (label = 1)
and then remove the node from the graph
and its adjacent edges.
- (II) Next, find another source node. We know
there be another source in the remaining graph
since the graph is still a DAG — removing
nodes and edges does not create cycles.
Label the next source as 2.
- (III) Repeat this procedure recursively.

How can we use Fact 4 to find a topological order?

The node that we label *first* in the topological sorting must have no incoming edges. Fact 4 guarantees that such a node exists.

Fact 5.

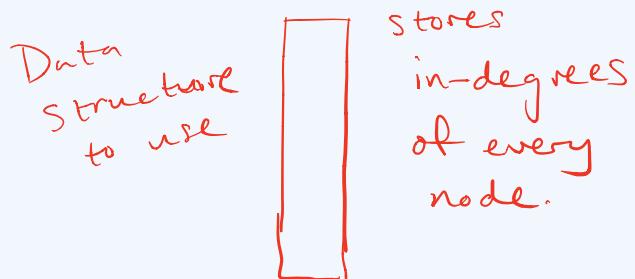
Let G' be the graph after a source node and its adjacent edges have been removed. Then G' is a DAG.

Proof: removing edges from G cannot yield a cycle!

This gives rise to a recursive algorithm for finding the topological order of a DAG. Its correctness can be shown by induction (use Facts 4, 5 to show induction step).

Algorithm for topological sorting

`TopologicalOrder(G)`



1. Find a source vertex s and order it first.
2. Delete s and its adjacent edges from G ; let G' be the new graph.
3. `TopologicalOrder(G')`
4. Append the order found after s .

Running time: $O(n^2)$. Can be improved to $O(n + m)$.

→ n calls to `TopologicalOrder(.)`

→ $O(n)$ per call

Running Time of Topological Order procedure

TopologicalOrder(G)

1. Find a source vertex s and order it first. $\rightarrow O(n)$
2. Delete s and its adjacent edges from G ; let G' be the new graph. $\rightarrow O(\deg(s))$
3. TopologicalOrder(G') \rightarrow
4. Append the order found after s . \rightarrow

We will introduce an array $\text{indegrees} =$

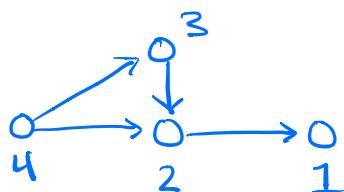
of size n , where :

$\text{indegrees}[i] = \text{indegree of node } i$

The array stores the indegree of each node.

This is useful for this since source nodes will have an indegree equal to 0.

Ex:



$\text{indegrees} =$

1	1
2	2
1	3
0	4

If a node v is such that $\text{indegrees}[v] = 0$, then that node is a source node. We use this array to identify a source node.

- How much time needed to initialize the array indegrees?

$$O(n+m)$$

Indeed, we need to scan the graph to determine how many in-neighbors every node has, and the graph has an adjacency list representation.

Therefore, we will need to scan the adjacency list representation, which will take $O(n+m)$ time. (Scanning list and updating indegrees).

- Step 1 running time:

How much time do you need to find a source vertex given the array indegrees (Step 1)?

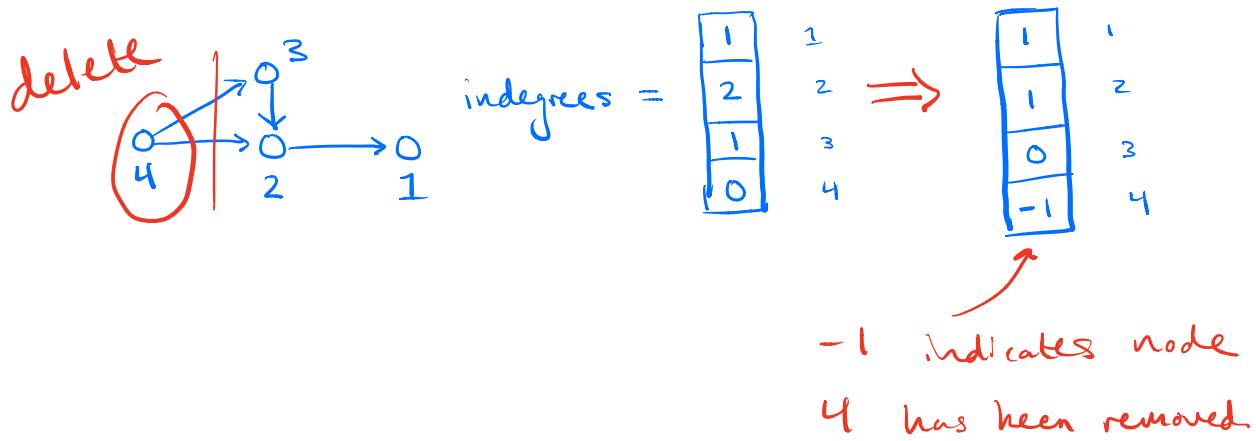
$$O(n)$$

We must scan the array indegrees to find a node with $\text{indegree} = 0$.

We are guaranteed such a node will exist since we are working with a DAG.

- Step 2 running time:

To remove a source node s and edges, all you need to do is update the array indegrees to delete the node. For example:



Then step 2 takes $O(\deg(s))$ time

for each source node.

Total running time:

$$n \cdot (O(n) + O(\deg(s)))$$

$$= O(n^2) + O\left(\sum_{s \in v} \deg(s)\right) =$$

$$= O(n^2) + O(m) = O(n^2)$$

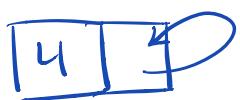
The bottleneck of this algorithm is
Step 1, identifying the source vertex.

Therefore, there is room for improvement
here.

To improve the algorithm, we will modify
• Step 2

Step 2: maintain a linked list with
source nodes. Every node v
with indegrees $[v] = 0$ will be added
to the linked list.

Sources



Initialize sources : $O(n)$

time

- Step 1 now takes $O(1)$ time because just need to extract head element.

* Maintain a linked list of sources.

Modified Algorithm :-

Step 1 : $O(1)$

Step 2 : $O(\deg(s))$

Total Running time :-

$$O(n+m) + n \cdot (O(1) + O(\deg(s))) + O(n)$$

$$= O(n+m) + O(n) + O(m) + O(n)$$

$$= O(n+m)$$

We will now cover yet another way
to do topological ordering / sorting :- ↴

Topological sorting via DFS

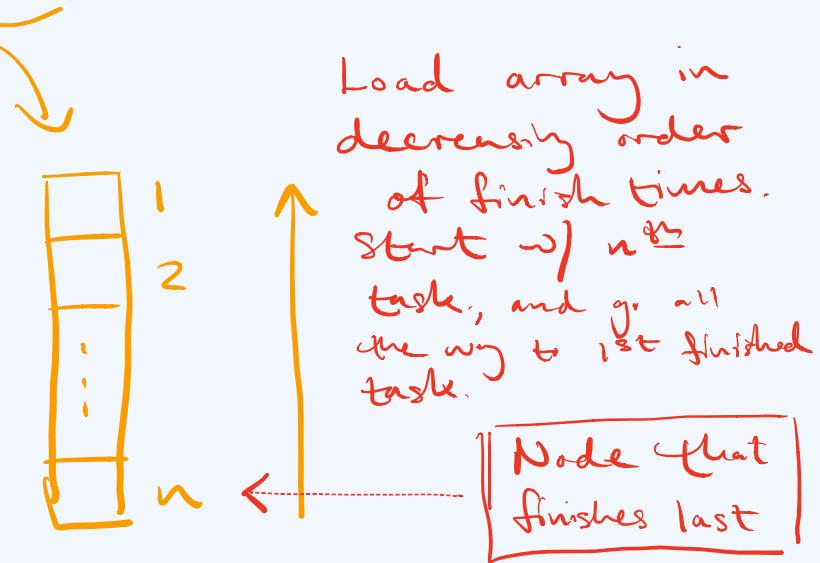
Let $G = (V, E)$ be a DAG.

Computing finish times requires $O(n+m)$ time.

- ▶ Run $\text{DFS}(G)$; compute *finish* times.
- ▶ Process the tasks in decreasing order of *finish* times.

Running time: $O(m + n)$

$F[i] = \text{node that finishes } i^{\text{th}} \text{ in DFS}$



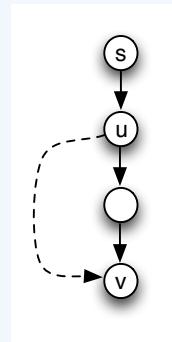
Intuition behind this algorithm

- ▶ The task v with the largest $finish$ has no incoming edges (if it had an incoming edge from some other task u , then u would have the largest $finish$). Hence v does not depend on any other task and it is safe to perform it first.
- ▶ The same reasoning shows that the task w with the second largest $finish$ has no incoming edges from any other task except (maybe) task v . Hence it is safe to perform w second.
- ▶ And so on and so forth.

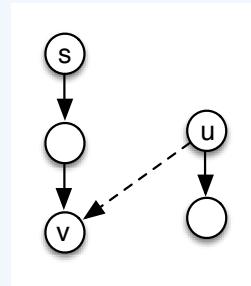
Formal proof of correctness

By Claim 4 there are no back edges in the DFS forest of a DAG. Thus every edge $(u, v) \in E$ is either

1. **forward/tree**: $start(u) < start(v) < finish(v) < finish(u)$



2. or **cross** edge: $finish(v) < start(u) < finish(u)$



Proof of correctness (cont'd)

Hence for every $(u, v) \in E$, $finish(v) < finish(u)$.

Consider a task v . All tasks u upon which v depends, that is, all tasks u such that there is an edge $(u, v) \in E$, satisfy $finish(v) < finish(u)$.

Since we are processing tasks in **decreasing** order of finish times, all tasks u upon which v depends have already been processed before we start processing v .

Let $G = (V, E)$ be a DAG.

- ▶ Run $\text{DFS}(G)$; compute *finish* times.
- ▶ Process the tasks in **decreasing** order of *finish* times.

Correctness: Let w be the task with largest finish time. We want to show that it is safe to perform w first.

Since w has the largest finish time it does not depend on any other task. The task we perform first should be a source task. w should not have any incoming edges.

Suppose to the contrary that $\exists (x, w) \in E$.

This means that x would also be discovered by DFS. There are 2 cases :

- (1) x is discovered by DFS before w .
- (2) x is discovered by DFS after w .

In Case (1), if x is discovered before

w , then x is an ancestor of w .



This implies that the finish time of x is larger than the finish time of w , a contradiction.

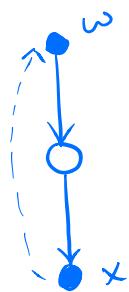
In Case ②, there are two sub-cases:

(I.) x discovered after w has finished $\Rightarrow (x, w)$ a cross edge.

This also implies that the finish time of x is greater than the finish time of w . Since for cross edges:

$$\text{start}(w) < \text{finish}(w) < \text{start}(x) < \text{finish}(x).$$

A contradiction.



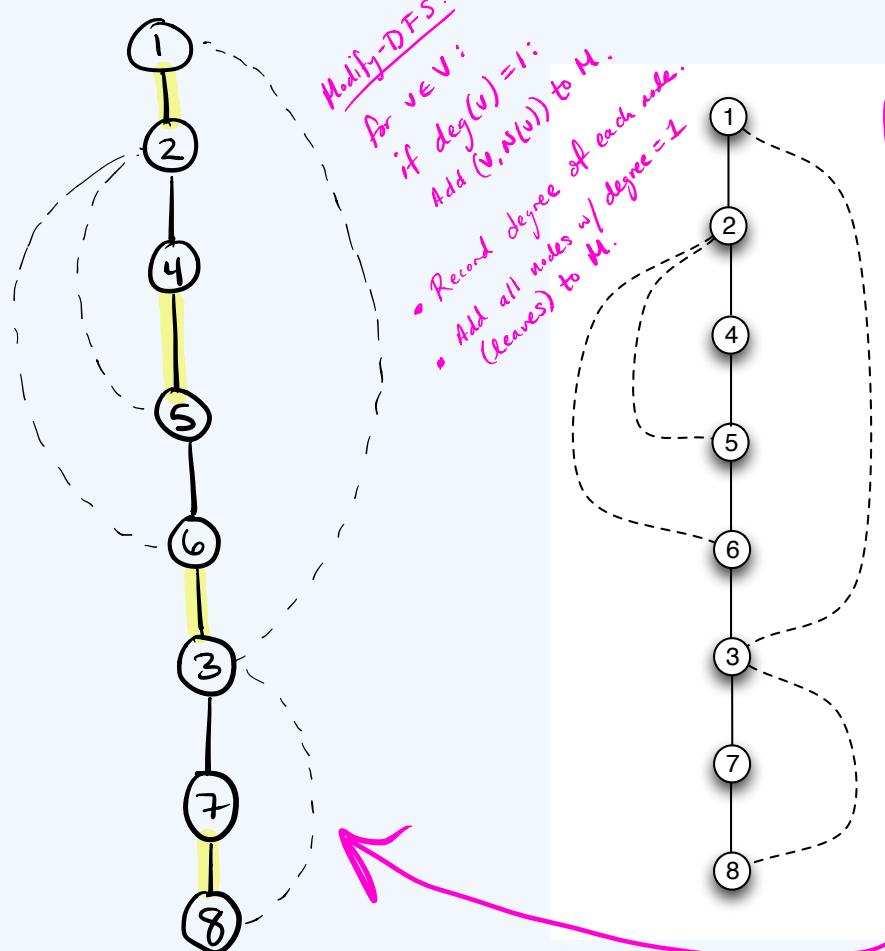
(II.) x discovered after w while w is still in the stack. Then, (x, w) is a back edge \Rightarrow the graph has a cycle, which contradicts G being a DAG. //

The following pages contain

SCRATCH WORK

for Homework problem
on Perfect Matchings

The DFS tree for the undirected graph G_1



if $\text{matched}[v] == \text{True}$:

- * In the DFS tree add every other edge to M .

DELETE WHEN DONE

Dashed edges belong to the graph but not to the DFS tree. Ties are broken by considering nodes by increasing index.

DFS($G = (V, E)$) Total time
 $O(n) \left\{ \begin{array}{l} \text{for } u \in V \text{ do} \\ \quad explored[u] = 0 \\ \text{end for} \\ \text{for } u \in V \text{ do} \\ \quad \text{if } explored[u] == 0 \text{ then Search}(u) \\ \quad \text{end if} \\ \text{end for} \end{array} \right.$
For fix
How many +
you can see
 $O(n) \left\{ \begin{array}{l} \text{Search}(u) \\ \rightarrow \text{previsit}(u) \rightarrow O(1) \\ \rightarrow explored[u] = 1 \rightarrow O(1) \quad * T(1) \\ \text{for } (u, v) \in E \text{ do} \\ \quad \text{if } explored[v] == 0 \text{ then Search}(v) \\ \quad \text{end if} \\ \text{end for} \\ \text{postvisit}(u) \rightarrow O(1) \end{array} \right.$
Let $T(u) =$
EX
SP
in
=
For fixed
 $T(u) =$

Running time for DFS if previsit, postv

for $u \in V$:

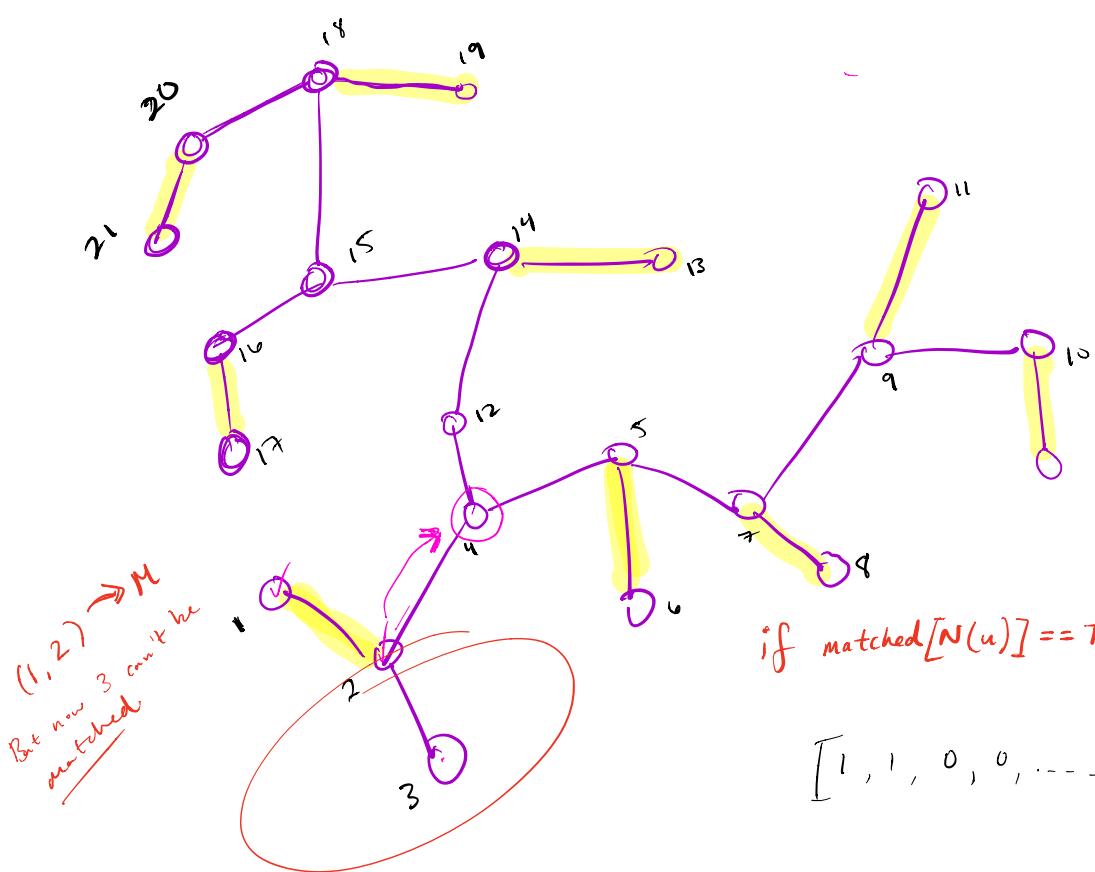
if matched $[u] == \text{False}$:

if $\deg(u) == 1$:

Add $(u, N(u))$ to M

for $u \in V$:
 if $\text{matched}[u] == \text{False}$:
 if $\deg(u) == 1$:

if matched[N(u)] == True :



if $\deg(u) == 1$ and $\text{matched}[u] == 0$ and
 $\text{matched}[N(u)] == 0$:

$\text{matched}[u] = 1$

$\text{matched}[N(u)] = 1$

else if $\deg(u) == 1$ and $\text{matched}[u] == 0$ and
 $\text{matched}[N(u)] == 1$:

$\text{matched}[u] = 0$

for $(u, v) \in E$:

if $\text{matched}[u] == 0$ AND $\text{matched}[v] == 0$

$\text{matched}[u] = 1$

$\text{matched}[v] = 1$

if $\deg(u) == 1$ and $\text{matched}[u] == 0$:

if $\text{matched}[N(u)] == 0$:

$\text{matched}[u] = 1$

$\text{matched}[N(u)] = 1$

else:

return 'False'

```
if sum(matched) == size(v) :  
    return 'True'
```

