

Analysis of Algorithms, I

CSOR W4231

Eleni Drinea
Computer Science Department

Columbia University

The dynamic programming principle; segmented least squares

Outline

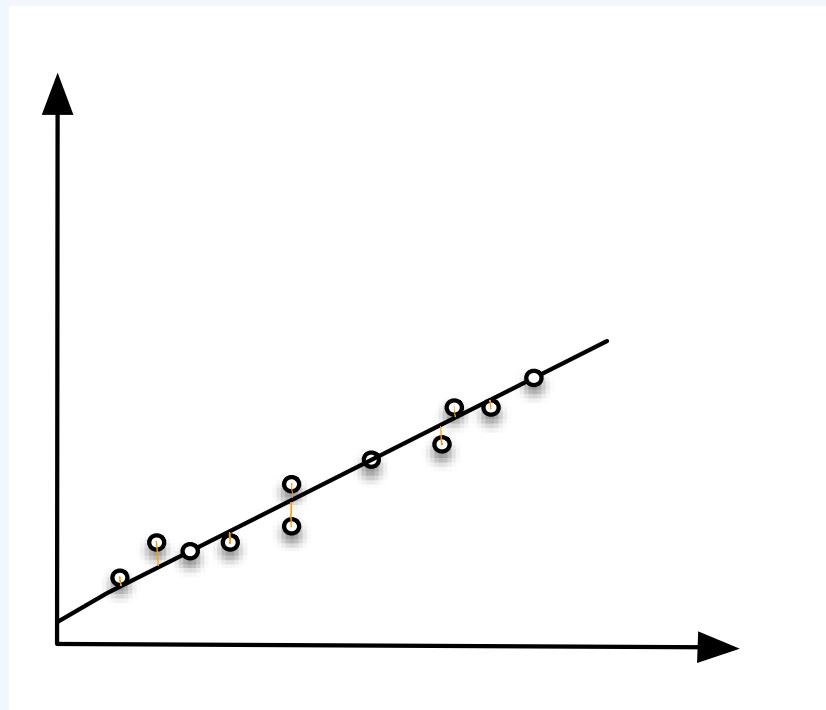
- 1 Segmented least squares
 - An exponential recursive algorithm
- 2 A Dynamic Programming (DP) solution
 - A quadratic iterative algorithm
 - Applying the DP principle

Today

- 1 Segmented least squares
 - An exponential recursive algorithm
- 2 A Dynamic Programming (DP) solution
 - A quadratic iterative algorithm
 - Applying the DP principle

Linear least squares fitting

A foundational problem in statistics: find a line of *best fit* through some data points.



Linear least squares fitting

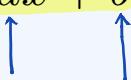
Input: a set P of n data points $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$;
we assume $x_1 < x_2 < \dots < x_n$.

Output: the line L defined as $y = ax + b$ that **minimizes** the
error

$\text{err}(L, P)$ is

a bivariate
function in a and b

$$\text{err}(L, P) = \sum_{i=1}^n (y_i - ax_i - b)^2$$



looking for an ' a ' and
' b ' such that it minimizes
the distance
to the line. (1)

minimizing the vertical distances of the data points to the
fitted line. Want to find an a ? b that minimize
the distances.

Linear least squares fitting: solution

Given a set P of data points, we can use calculus to show that the line L given by $y = ax + b$ that minimizes

Given a and b
 $\text{err}(L, P)$ can be
computed in $O(n)$.

$$\text{err}(L, P) = \sum_{i=1}^n (y_i - ax_i - b)^2 \quad (2)$$

satisfies

$$O(n) \quad \curvearrowright O(n)$$

$$a = \frac{n \sum_i x_i y_i - (\sum_i x_i)(\sum_i y_i)}{n \sum_i x_i^2 - (\sum_i x_i)^2} \quad \curvearrowright O(n) \quad (3)$$

$$b = \frac{\sum_i y_i - a \sum_i x_i}{n} \quad \longrightarrow O(n) \quad (4)$$

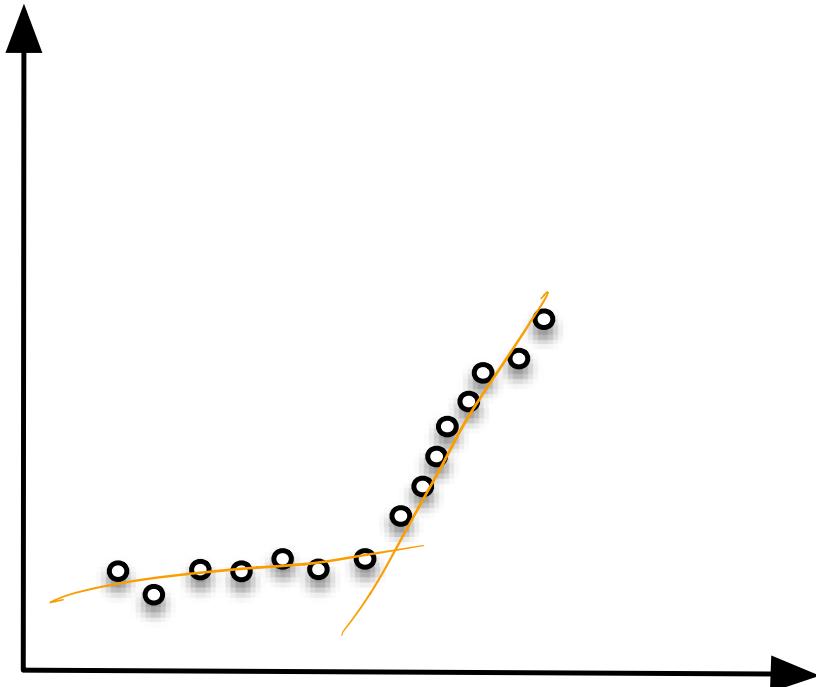
Given 'a' you
can compute 'b'
in $O(n)$ time

How fast can we compute a, b ? \rightarrow linear time.

$O(n)$ for 'a': Use a single for-loop to compute $\sum x_i y_i$, $\sum x_i$, $\sum y_i$, and $\sum x_i^2$. Can compute b in $O(n)$ time too.

What if the data changes direction?

• Could fit 2 lines to the data instead of just 1 line to minimize the error w.r.t. the data.



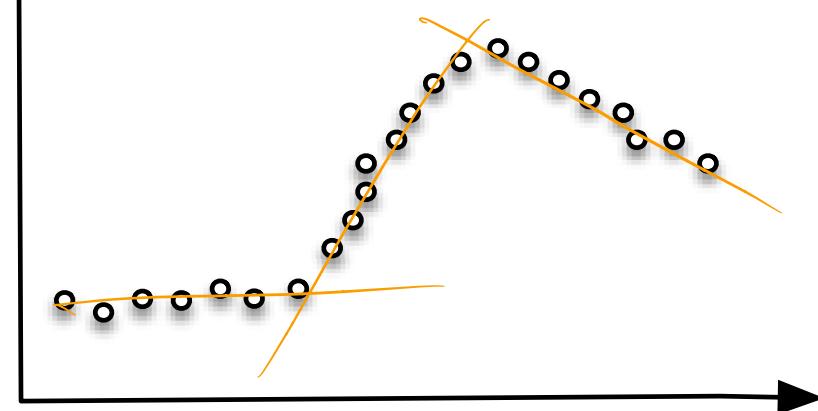
What if the data changes direction more than once?

Here, 2 lines will not suffice.

- Will need 3 lines.

However, don't want to look at data and try to guess

Want to introduce the fewest lines possible to minimize the error.



- Want an algorithm that detects change in direction for any input, and therefore introduces new lines as needed to fit the data well.

The # of lines needed to fit the data and hardwire that into your code. We want an algorithm that automatically detects change in direction.

What if the data changes direction more than once?

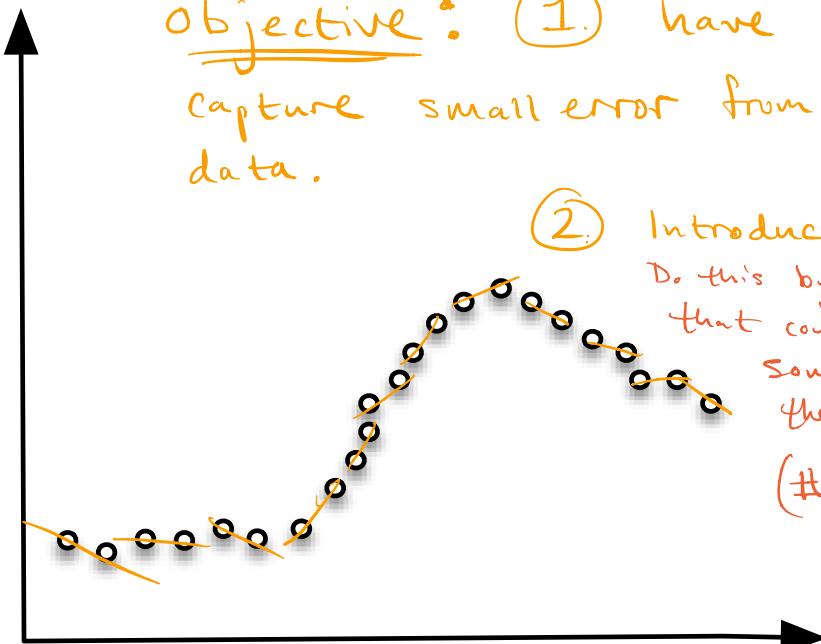
Real Goal: Introduce FEW lines, while keeping the OVERALL error small.

Objective: ① have a term to capture small error from fitting lines to data.

② Introduce few lines.
Do this by incorporating a term that counts # of lines times some cost for introducing the line:

$$(\# \text{ lines}) \cdot C$$

regularization
parameter



Don't want to make objective super hard.

* We have 3rd implicit goal which is to come up w/ an optimization problem that we can actually solve efficiently.

Our objective will be a simple sum of the error of fitting lines to the data plus a cost for introducing a new line times the number of lines. :

$$\text{obj} = \{\text{error for fitting lines to data}\} + \{C \cdot \# \text{ lines}\}$$

where C is the cost for introducing a line.
(cost per 1 line).

This captures the trade-off between accuracy and simplicity of the model.

$$\text{obj} = \underbrace{\{\text{error for fitting lines to data}\}}_{\text{keeping the error small increases the # of lines}} + \underbrace{\{C \cdot \# \text{ lines}\}}_{\text{keeping overall cost low increases the error.}}$$

$$\downarrow \text{Error} \Rightarrow \uparrow \{C \cdot \# \text{ of lines}\}$$

$$\downarrow \# \text{ lines} \Rightarrow \uparrow \text{Error}$$

Keeping C small \Rightarrow Care more about accuracy than simplicity.
(or smaller error)

Large C \Rightarrow Care more about simplicity of the model because adding lines causes the objective to grow a lot.

By fine tuning C , you can err towards simplicity or increased accuracy for the model.

C is a regularization parameter.

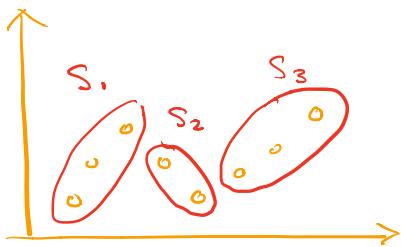
- if you care more about accuracy, keep C small.
- if you care more about simplicity, keep C large.

What we're doing :

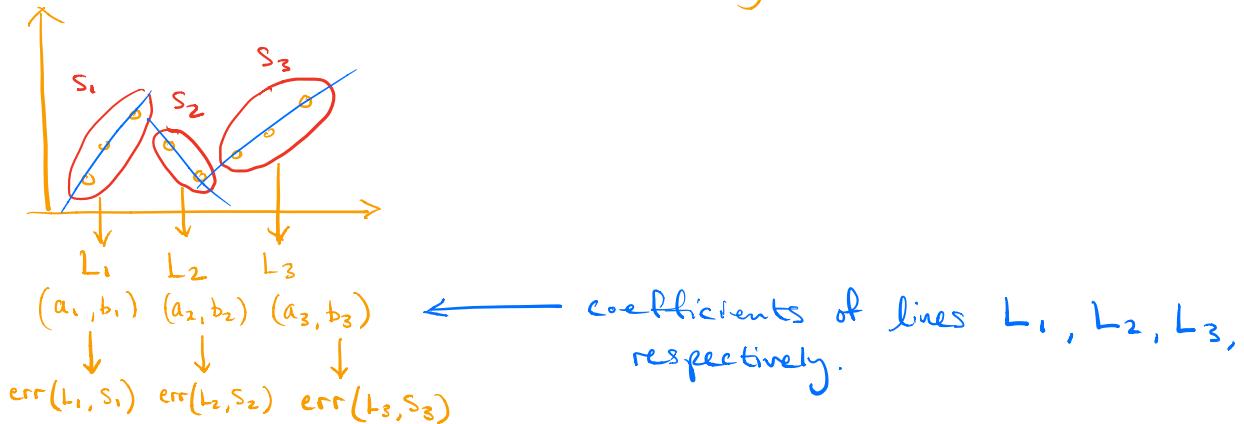
Input data that change direction:



1. Partition input into segments of consecutive data points



2. Fit lines of best fit to each segment to minimize the total error of fitting lines



Segment S_i : contiguous set of data points
 $\{P_1, P_2, \dots, P_j\}$

Partition or Segmentation of the input P :

a partitioning of the input into m

segments: S_1, S_2, \dots, S_m (call this Partition A)

Cost of Partition A :

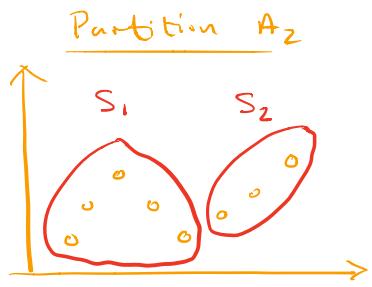
$$\sum_{i=1}^m \underbrace{\text{err}(L_i, S_i)}_{\substack{\text{error of line } L_i \\ \text{in segment } S_i}} + \underbrace{m \cdot C}_{\substack{\text{total cost of} \\ \# \text{ of lines} \\ \text{used.}}} \quad (m = \# \text{ of segments or lines})$$

model error

$m \equiv \# \text{ of segments (and thus the \# of lines)}$.



$$\text{cost}(A_1) = \sum_{i=1}^3 \text{err}(L_i, S_i) + 3C$$



$$\text{cost}(A_2) = \sum_{i=1}^2 \text{err}(L_i, S_i) + 2C$$



Specifying a Partition A : There are a number of ways. A few ways include:

- By segment of pts: $\{P_1, \dots, P_k\}$, $\{P_{k+1}, \dots, P_\ell\}$, $\{P_{\ell+1}, \dots, P_n\}$
- By starting points of each segment:
 $\{P_1, P_6, P_{11}\}$
- By the last points of each segment:
 $\{P_5, P_{10}, P_{20}\} \xrightarrow{\text{equiv.}} \{P_5, P_{10}\}$ equivalent since it's clear P_{20} is last.
- # of segments & the end pts: $\{3, P_5, P_{10}, P_{20}\}$

How to detect change in the data

- ▶ Any single line would have large error.
- ▶ **Idea 1:** hardcode number of lines to 2 (or some *fixed m*).
 - ▶ Fails for the dataset on the last slide.
- ▶ **Idea 2:** pass an *arbitrary set* of lines through the points and seek the set of lines that minimizes the error.
 - ▶ Trivial solution: have a different line pass through each pair of consecutive points in P .
- ▶ **Idea 3:** fit the points well, using as few lines as possible.
 - ▶ Trade-off between complexity and error of the model

Formalizing the problem

Input: data set $P = \{p_1, \dots, p_n\}$ of points on the plane.

- ▶ A **segment** $S = \{p_i, p_{i+1}, \dots, p_j\}$ is a contiguous subset of the input.
- ▶ Let \mathcal{A} be a partition of P into $m_{\mathcal{A}}$ segments $S_1, S_2, \dots, S_{m_{\mathcal{A}}}$.
For every segment S_k , use (2), (3), (4) to compute a line L_k that minimizes $err(L_k, S_k)$.
- ▶ Let $C > 0$ be a fixed multiplier. The **cost** of partition \mathcal{A} is

$$\sum_{S_k \in \mathcal{A}} err(L_k, S_k) + m_{\mathcal{A}} \cdot C$$

Segmented least squares

This problem is an instance of change detection in data mining and statistics.

Input: A set P of n data points $p_i = (x_i, y_i)$ as before.

Output: A segmentation $\mathcal{A}^* = \{S_1, S_2, \dots, S_{m_{\mathcal{A}^*}}\}$ of P whose

cost

$$\sum_{S_k \in \mathcal{A}^*} err(L_k, S_k) + m_{\mathcal{A}^*} C$$

is minimum.

# of segments		# partitions w/ m segments
$m=1$	$\frac{1}{0} \quad \frac{2}{0} \quad 0 \dots \frac{n}{0}$	$1 = \binom{n-1}{0}$
$m=2$	$\frac{1}{0} \quad \frac{2}{0} \quad \quad 0 \dots \frac{n-1}{0} \quad \quad \frac{n}{0}$ ↑ ↑ ↑ Stopping Points (Recall, pts must be consecutive) * $n-1$ choices to end your first segment	$n-1$ $= \binom{n-1}{1}$
$m=3$	$\frac{1}{0} \quad \frac{2}{0} \quad 0 \dots \frac{n-1}{0} \quad \frac{n}{0}$	$\binom{n-1}{2}$
	• The last point must be n , and need to find 2 places to end the first two segments, and there are $\binom{n-1}{2}$ choices of this	
⋮	⋮	⋮
$m=k$	Need to find $k-1$ places to end the first $k-1$ segments.	$\binom{n-1}{k-1}$
⋮	⋮	⋮
$m=n$	$\frac{1}{0} \quad \frac{2}{0} \quad \quad 0 \quad \dots \quad \frac{n-1}{0} \quad \quad \frac{n}{0} \quad $	1 $= \binom{n-1}{n-1}$
	Each point is its own segment.	

$$\sum_{k=0}^{n-1} \binom{n-1}{k} = 2^{n-1}$$

Total # of
possible partitions

$$= \underline{\Omega}(2^n) = \underline{\Omega}(2^n)$$

Recall:

$$\sum_{i=0}^n \binom{n}{i} \cdot a^i \cdot b^{n-i} = (a+b)^n$$

For $a=b=1$:

$$\sum \binom{n}{i} \cdot a^i \cdot b^{n-i} = \sum_{i=0}^n \binom{n}{i} = 2^n$$

Brute force
algorithm requires
exponential time

Issue w/ this count:

- No reason to have partitions with $m > \frac{n}{2}$.

$\frac{n}{2}$ segments means every segment has two data points, and there is only one possible partitioning of n pts into $\frac{n}{2}$ segments of consecutive data points.

⇒ The max # of lines to use is $\frac{n}{2}$ and there is only one possible partition w/ $\frac{n}{2}$ segments. 

The count however is much larger
for $n/2$, but should only be 1.

Need to get more appropriate bound.

Closer to: $\underline{\Omega}(2^{n/2}) = \Omega(2^{n/2})$

(underestimate for the true # of partitions)

- Brute force algorithm requires exponential time.
- Lowerbound of brute force approach is still exponential \Rightarrow Inefficient.

NOTE: If you want to show an algorithm is inefficient, you need to show that the lowerbound is exponential (or inefficient)

A brute force approach

We can find the optimal partition (that is, the one incurring the minimum cost) by exhaustive search.

- ▶ Enumerate every possible partition (segmentation) and compute its cost.
- ▶ Output the one that incurs the minimum cost.

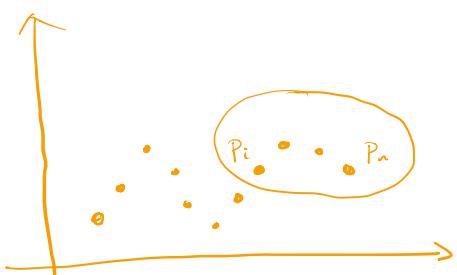
△ $\Omega(2^n)$ partitions

$\overbrace{\quad}^{\infty}$ inefficient

Aim: Recursive approach / algorithm.

- Extract/Remove the last segment of the optimal segmentation.
- Continue recursively.

→ Let A^* be the optimal segmentation.



- Extract the last segment S .
- We know that $p_n \in S$.
- Do not know: where S starts.
- We know that S starts at some point p_i , where p_i could be p_1 or p_2 or p_3 or ... or p_n ($1 \leq i \leq n$)

Assume that S starts at p_i :

Let $\text{OPT}(n) = \min \text{ cost of segmenting } p_1, \dots, p_n$

- i.e., the minimum cost of segmenting the entire input.

Can we express this recursively?

$$\text{OPT}(n) = C + \underset{\uparrow}{\text{err}}(L, \{p_i, \dots, p_n\}) + \underset{\uparrow}{\text{OPT}(i-1)}$$

$$OPT(n) = \min_{1 \leq i \leq n} \left\{ C_i + err(L, \{P_i, \dots, P_n\}) + OPT(i-1) \right\}$$

cost/contrtribution of last segment

min cost to segment
 P_i, \dots, P_{i-1}

We do not know π though - need to find it. :

Your p_i should be such that it minimizes the total cost. That is, we minimize over p_i (the starting points) :

$$OPT(n) = \min_{1 \leq i \leq n} \left\{ \underbrace{C + \text{err}(L, \{P_i, \dots, P_n\})}_{\begin{array}{l} \text{still need} \\ \text{to segment rest} \end{array}} + OPT(i-1) \right\}$$

minimize over every possible choice for p_i \uparrow
 We try every possible p_i and find the one that minimizes
 the expression.

Boundary condition :

$$\text{OPT}(0) = 0$$

This is the recursive algorithm.

To apply recursively to remove the last segment, define:

- $e_{ij} \triangleq \text{err}(L, \{P_i, \dots, P_j\})$

e_{ij} ≡ error of fitting a line to the segment that starts at P_i

PRECOMPUTE
 $\frac{1}{2}$
 STORE
 e_{ij} in $O(n^2)$

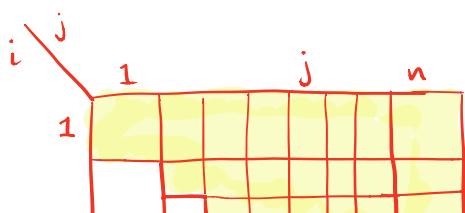
$$\text{OPT}(j) = \min_{1 \leq i \leq j} \left\{ e_{ij} + C + \text{OPT}(i-1) \right\}$$

\curvearrowright Can access e_{ij} in $O(1)$ time

- We pre-compute the e_{ij} values and store them in a matrix, so that whenever I need to look at them I can do so in constant time.

$$1 \leq i \leq j \leq n \Rightarrow \Theta(n^2) \text{ } e_{ij} \text{ terms}$$

- if $i < j$ we have $\binom{n}{2}$ e_{ij} terms.
- if $i \leq j$ we have $\binom{n}{2} + n$ e_{ij} terms.



- Segments to consider:
- $\{P_1\}, \{P_1, P_2\}, \dots, \{P_1, \dots, P_n\}$
 - $\{P_2\}, \dots, \{P_2, \dots, P_n\}$
 - . . .

$$E = \begin{matrix} i \\ \vdots \\ n \end{matrix} \quad \begin{matrix} j \\ \vdots \\ n \end{matrix}$$

e_{ij} in time: $O(j-i)$
 $= O(n)$

E in time: $O(n^3)$

Interested in filling in the upper triangular region of this matrix, including the diagonal entries. There are a set - n^2 entries in total in the matrix E.

- Time to fill in $e_{ij} = O(j-i) = O(n)$

\Rightarrow Total time to fill in E = $O(n^3)$

since there are n^2 entries.

and each entry takes $O(n)$ time.

→ Can be improved to $O(n^2)$

$O(n)$ time
 to compute a_i
 and b_i for
 $L_i = a_i x + b_i$
 Then
 $\text{err}(L_i, S_i)$
 takes
 $O(n)$
 time.

PRECOMPUTE \in STORE E

- time : $O(n^2)$
 - space : $O(n^2)$

Running Time :

$$OPT(n) = \min_{1 \leq i \leq n} \{ e_{ij} + C + OPT(i-1) \}$$

Let $T(n) \equiv$ time to compute $\text{OPT}(n)$

$$T(n) = T(n-1) + T(n-2) + T(n-3) + \dots$$

$\Omega \in \mathcal{N}$

$$\dots + \dots T(0) + \underbrace{T(i-1) + \underline{\Omega}(1)}$$

Need to compute $\text{OPT}(i)$ for
 $1 \leq i \leq n$:

$$\text{OPT}(n) = \min \left\{ \begin{array}{l} c_{n,n} + C + \text{OPT}(n-1), \\ c_{n,n-1} + C + \text{OPT}(n-2), \\ c_{n,n-2} + C + \text{OPT}(n-3), \\ \dots \end{array} \right\}$$

compute $c_{ij} + C + \text{OPT}(i-1)$
for a fixed i
And $T(i-1) + \underline{\Omega}(1) = \underline{\Omega}(1)$

$$\begin{aligned} \text{BAD!} &= T(n-1) + T(n-2) + \dots + T(0) + \underline{\Omega}(1) \\ &\geq T(n-1) + T(n-2) \end{aligned}$$

All you
need to
prove a
negative
result
is to show
it has a
bad lowerbound.

Prove that
it is super
polynomial?

$$T(n) \geq T(n-1) + T(n-2)$$

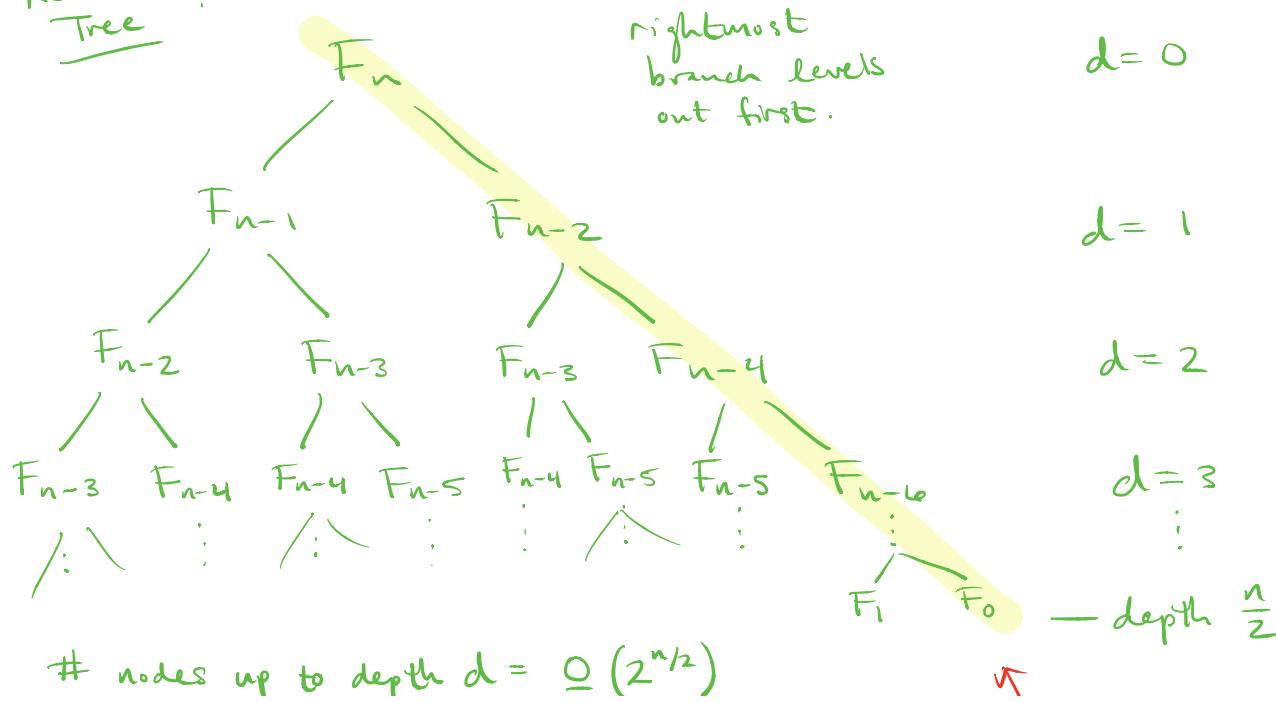
Look familiar?

Fibonacci Numbers!

$$F_n = F_{n-1} + F_{n-2}, \quad \begin{cases} F_0 = 0 \\ F_1 = 1 \end{cases}$$

Recursion:

Tree



- At depth $\frac{n}{2}$, the right-most branch bottoms out (first one to bottom out)
 - There are $2^{n/2}$ leaves in a complete binary tree.

\Rightarrow If you sum up all the work spent in this tree to get F_n , you are summing up at least $2^{n/2} \cdot c$

other branches
 not done, but
 this already
 requires $\Omega(2^{n/2})$

Therefore, $F_n \geq 2^{n/2}$

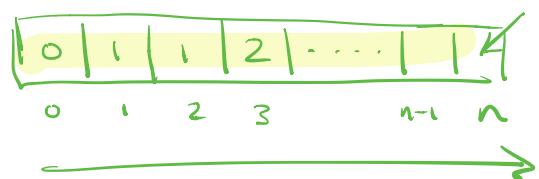
and thus $T(n) = \Omega(2^{n/2})$ ← Super Polynomial time
 (Inefficient!)

NOTE: Computing the n^{th} Fibonacci # does NOT require exponential time.

Create an array F with $F[0] = 0$ and

$$F[1] = 1 \quad ; \quad F = \boxed{0 \mid 1 \mid \text{ } \mid \text{ } \mid \text{ } \mid \cdots \mid \text{ } \mid n}$$

to store the Fibb. #'s, compute F_n from the bottom up. :



Time to compute F_n this way = $O(n)$

(assuming it requires $O(1)$ time to combine 2 #'s to compute a single value)

- Doing it this way we are only computing one branch of the tree by storing in the array, solutions to subproblems that were already previously computed, so that you can look them up in constant time in the future. We can do this because $F_n = F_{n-1} + F_{n-2}$ has overlapping subproblems.

$$\text{OPT}(n) = \min_{1 \leq i \leq n} \{ e_{ij} + C + \text{OPT}(i-1) \}$$

also has overlapping subproblems, so we can compute this faster using a similar approach described above, using a bottom-up approach. (Generally, this is how we will approach dynamic programming problems.) We need space to store solutions.

$$OPT(n) = \min_{1 \leq i \leq n} \{ e_{ij} + c + OPT(i-1) \}$$

Introduce an array M to store solutions.

$M[i] = OPT(i)$ with $n+1$ entries.

M	<table border="1" style="border-collapse: collapse; width: 100px; margin-bottom: 5px;"> <tr> <td style="width: 20px; height: 20px;"></td><td style="width: 20px; height: 20px;"></td><td style="width: 20px; height: 20px;"></td><td style="width: 20px; height: 20px;"></td></tr> <tr> <td style="text-align: center;">0</td><td style="text-align: center;">1</td><td style="text-align: center;">2</td><td style="text-align: center;">...</td></tr> </table>					0	1	2	...	$\overbrace{\hspace{100px}}$	n	$(OPT(0) = 0)$
0	1	2	...									

- Fill it in from the bottom up, using the recurrence.

$$\begin{aligned}
 OPT(1) &= \cancel{e_n}^0 + c + OPT(0) \\
 &= c \quad e_n = 0 \text{ since fitting line to 1 point.}
 \end{aligned}$$

$$\begin{aligned}
 OPT(2) &= \min_{1 \leq i \leq 2} \left\{ \begin{array}{l} e_{12} + c + OPT(0), \\ e_{22} + c + OPT(1) \end{array} \right\} \\
 &= \min_{1 \leq i \leq 2} \{ c, 2c \} = c
 \end{aligned}$$

$$OPT(3) = \min_{1 \leq i \leq 3} \left\{ \begin{array}{l} e_{13} + c + OPT(0), \\ e_{23} + c + OPT(1), \\ e_{33} + c + OPT(2) \end{array} \right\}$$

$$= \min_{1 \leq i \leq 3} \{ e_{i,3} + c, 2c, 3c \}$$

$$= X$$

$$M = \boxed{0 | c | c | X | \dots | }$$

0 1 2 3 n



How can we prove correctness?

Induction *(Do on our own
as an exercise)*

- EASY TO COMPUTE RECURRENCE

$$F_n = F_{n-1} + F_{n-2}$$

- STORE & RE-USE solutions to subproblems



- ITERATIVE BOTTOM-UP COMPUTATIONS

$$F[0] = 0$$

$$F[1] = 1$$

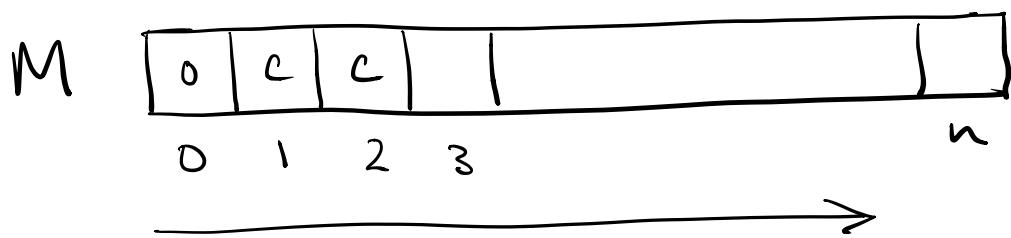
$$F[i] = F[i-1] + F[i-2] \quad \text{for } i > 2$$

- We want : F_n
 - Time per subproblem : $O(1)$
 - # subproblems : n
 - Total time : $O(n)$
 - Space : $O(n)$

$$OPT(j) = \min_{1 \leq i \leq j} \{e_{i,j} + C + OPT(i-1)\}$$

$$M[j] = OPT(j)$$

Computed in
a bottom-up
fashion.



$$M[0] = OPT(0) = 0$$

$$M[1] = OPT(1) = c$$

$$\begin{aligned} M[2] &= OPT(2) = \min \left\{ \begin{array}{l} e_{12}^0 + c + OPT(0), \\ e_{22}^0 + c + \underbrace{OPT(1)}_{=c} \end{array} \right\} \\ &= \min \{c, 2c\} = c \end{aligned}$$

$$M[3] = OPT(3) = \min_2 \left\{ \begin{array}{l} e_{13}^0 + c + OPT(0) = \\ e_{23}^0 + c + OPT(1) = 2c \\ e_{33}^0 + c + OPT(2) = 2c \end{array} \right.$$

$$M[4] = OPT(4) =$$

$$OPT(j) = \min_{1 \leq i \leq j} \{e_{i,j} + C + OPT(i-1)\}$$

Time for $OPT(j)$: $O(j)$
 $= O(n)$

- We want : $M[n] = OPT(n)$
- Time per subproblem : $O(n)$ ← Because minimizing over j entries.
- # subproblems : n
- Total time : $O(n^2)$
- Total space : $\underbrace{O(n)}_{\text{for dynamic portion of storage } (M)} + \underbrace{O(n^2)}_{\text{for } e_{i,j} \text{ matrix}}$

A crucial observation regarding the last data point

Consider the last point p_n in the data set.

- ▶ p_n belongs to a single segment in the **optimal** partition.
- ▶ That segment starts at an earlier point p_i , for some $1 \leq i \leq n$.

This suggests a **recursive** solution: **if** we knew where the last segment starts, then we could remove it and recursively solve the problem on the remaining points $\{p_1, \dots, p_{i-1}\}$.

A recursive approach

- ▶ Let $OPT(j) = \text{minimum cost of a partition of the points } p_1, \dots, p_j.$
- ▶ Then, if the last segment of the optimal partition is $\{p_i, \dots, p_n\}$, the cost of the optimal solution is

$$OPT(n) = err(L, \{p_i, \dots, p_n\}) + C + OPT(i - 1).$$

- ▶ But we don't know where the last segment starts! *How do we find the point p_i ?*
- ▶ Set

$$OPT(n) = \min_{1 \leq i \leq n} \left\{ err(L, \{p_i, \dots, p_n\}) + C + OPT(i - 1) \right\}.$$

A recurrence for the optimal solution

Notation: let $e_{i,j} = \text{err}(L, \{p_i, \dots, p_j\})$, for $1 \leq i \leq j \leq n$.

Then

$$OPT(n) = \min_{1 \leq i \leq n} \left\{ e_{i,n} + C + OPT(i-1) \right\}.$$

If we apply the above expression recursively to remove the last segment, we obtain the recurrence

$$OPT(j) = \min_{1 \leq i \leq j} \left\{ e_{i,j} + C + OPT(i-1) \right\} \quad (5)$$

Time to compute $OPT(j) \rightarrow \Theta(j) = O(n)$ (*and there are $n+1$ subproblems*)

Remark 1.

1. We can precompute and store all $e_{i,j}$ using equations (2), (3), (4) in $O(n^3)$ time. *Can be improved to $O(n^2)$.*
2. The natural recursive algorithm arising from recurrence (5) is **not** efficient (think about its recursion tree!).

\Rightarrow Time to compute M (entire algorithm) : $n \cdot O(n) = O(n^2)$

Exponential-time recursion

$$T(n) \geq F_n \geq 2^{n/2}$$

Notation: $T(n)$ = time to compute $OPT(n)$, that is, the cost of the optimal partition for n points.

Then

$$T(n) \geq T(n-1) + T(n-2).$$

- ▶ Can show that $T(n) \geq F_n$, the n -th Fibonacci number (by strong induction on n).
 - ▶ From optional problem 6a in Homework 1, $F_n = \Omega(2^{n/2})$.
 - ▶ Hence $T(n) = \Omega(2^{n/2})$.
- ⇒ The recursive algorithm requires $\Omega(2^{n/2})$ time.

Today

- 1 Segmented least squares
 - An exponential recursive algorithm
- 2 A Dynamic Programming (DP) solution
 - A quadratic iterative algorithm
 - Applying the DP principle

Are we really that far from an efficient solution?

Recall Fibonacci problem from HW1: exponential recursive algorithm, **polynomial** iterative solution

How?

1. **Overlapping subproblems:** spectacular redundancy in computations of recursion tree
2. **Easy-to-compute recurrence** for combining the smaller subproblems: $F_n = F_{n-1} + F_{n-2}$
3. **Iterative, bottom-up computations:** we computed and stored the subproblems from smallest (F_0, F_1) to largest (F_n), iteratively.
4. **Small number of subproblems:** only solved $n - 1$ subproblems.

Elements of DP in segmented least squares

Our problem exhibits similar properties.

1. Overlapping subproblems
2. Easy-to-compute recurrence for combining optimal solutions to smaller subproblems into the optimal solution of a larger subproblem (once smaller subproblems have been solved)
3. Iterative, bottom-up computations: compute the subproblems from smallest (0 points) to largest (n points), iteratively.
4. Small number of subproblems: we only need to solve n subproblems.

A dynamic programming approach

$$OPT(j) = \min_{1 \leq i \leq j} \left\{ e_{i,j} + C + OPT(i - 1) \right\}$$

- ▶ The optimal solution to the subproblem on p_1, \dots, p_j contains optimal solutions to smaller subproblems.
- ▶ Recurrence 5 provides an **ordering** of the subproblems from smaller to larger, with the subproblem of size 0 being the smallest and the subproblem of size n the largest.
- ⇒ There are $n + 1$ subproblems in total. Solving the j -th subproblem requires $\Theta(j) = O(n)$ time.
- ⇒ The overall running time is $O(n^2)$.
- ▶ Boundary conditions: $OPT(0) = 0$.
- ▶ Segment p_k, \dots, p_j appears in the optimal solution only if the minimum in the expression above is achieved for $i = k$.

An iterative algorithm for segmented least squares

This algorithm computes the optimal cost.

Let M be an array of n entries such that

$M[i] = \text{cost of optimal partition of the first } i \text{ data points}$

SegmentedLS(n, P)

$$M[0] = 0$$

for all pairs $i \leq j$ do

 Compute $e_{i,j}$ for segment p_i, \dots, p_j using (2), (3), (4)

end for

for $j = 1$ to n do

$$M[j] = \min_{1 \leq i \leq j} \{e_{i,j} + C + M[i - 1]\}$$

end for

Return $M[n]$

Precompute
and
store the
 e_{ij} 's

The heart of the
Dynamic Programming
algorithm.

Running time: time required to fill in dynamic programming array M is $O(n^3) + O(n^2)$. Can be brought down to $O(n^2)$.

We have $n+1$ subproblems, each requires $O(n)$ time.

Reconstructing an optimal segmentation

$O(n^2)$ is tight

Since this is an optimization problem, we also want to output an optimal solution (optimal segmentation here).

We can reconstruct the optimal partition recursively, using array M and error matrix e .

OPTSegmentation(j) (initially call it with argument n)

if ($j == 0$) then return

else

{ Find $1 \leq i \leq j$ such that $M[j] = e_{i,j} + C + M[i - 1]$

OPTSegmentation($i - 1$)

Output segment $\{p_i, \dots, p_j\}$

end if

► Initial call: OPTSegmentation(n)

► Running time?

• At most $O(n)$ recursive calls

• Time spent per recursive call : $O(j) = O(n)$

Finds the point where the last segmentation in the optimal segment up to point j starts.

up to $n/2$ segments

$O(n^2)$ time

Extracts each segment.

Reconstructing an optimal segmentation

We can reconstruct the optimal partition **recursively**, using array M and error matrix e .

```
OPTSegmentation(j)
```

```
  if ( $j == 0$ ) then return
```

```
  else
```

```
    Find  $1 \leq i \leq j$  such that  $M[j] = e_{i,j} + C + M[i - 1]$ 
```

```
    OPTSegmentation( $i - 1$ )
```

```
    Output segment  $\{p_i, \dots, p_j\}$ 
```

```
  end if
```

$p_i \leftarrow p_n$

$O(|S_i|)$

- ▶ Initial call: `OPTSegmentation(n)`
- ▶ *Running time?*

- We are usually stingier with space than we are with time, unless it helps improve the running time of the algorithm.
- To give the optimal solution, we typically use a recursive algorithm.



~~1 \rightarrow i~~ $P_i \leftarrow P_n$

- ~~1 \rightarrow i~~ might spend a long time to find the i where the last segment starts.

— Start from the opposite direction.
Instead of starting 'from 1 to i ', start from point n and go down looking for point i , then you are guaranteed to spend time proportional to the length of the segment to find its starting point. Do this for every segment, so total time spent is proportional to the sum of the lengths of all the segments, which is n .

This modification guarantees linear time.

Obtaining efficient algorithms using DP

1. Optimal substructure: the optimal solution to the problem contains optimal solutions to the subproblems.
2. A recurrence for the overall optimal solution in terms of optimal solutions to appropriate subproblems. The recurrence should provide a natural ordering of the subproblems from smaller to larger and require polynomial work for combining solutions to the subproblems.
3. Iterative, bottom-up computation of subproblems, from smaller to larger.
4. Small number of subproblems (polynomial in n).

Time for DP algorithms : $O(\# \text{ subproblems} \cdot \text{time per subproblem})$

For DP you need:

1. Recurrence of subproblems

(- need to define subproblems clearly)
(Ex: In this problem, $\text{OPT}(n) \equiv \min \text{ cost to segment end pts.}$)

2. Boundary conditions.

Ex: In this problem $\text{OPT}(0) = 0$

3. Order of subproblems from smaller to larger.

4. Iterative computation.

5. Space: $O(\# \text{ subproblems})$

6. Time: $O(\# \text{ subproblems} \cdot \frac{\text{time per}}{\text{subproblem}})$

* Solution to the original problem appear somewhere in the DP table, or somehow be derived from subproblems.

HW #3 deals w/ the latter form: Need to derive the solution to the overall problem by looking at some entries in your DP table - you will not find a single entry to give the solution to the

Dynamic programming vs Divide & Conquer

overall problem.

- ▶ They both combine solutions to subproblems to generate the overall solution.
- ▶ However, divide and conquer starts with a large problem and divides it into small pieces.
- ▶ While dynamic programming works from the bottom up, solving the smallest subproblems first and building optimal solutions to steadily larger problems.